

Package ‘rpbnb’

September 7, 2026

Type Package

Title Random-Parameter Bivariate Negative Binomial Regression

Version 0.4.8

Author Zhenyu Wang [aut, cre]

Maintainer Zhenyu Wang <wonstran@hotmail.com>

Description Maximum-likelihood estimation of bivariate negative binomial (NB2) regression models with Famoye/Sarmanov or discrete-copula (Frank, Gaussian, or Clayton) dependence, and maximum-simulated-likelihood estimation of a bivariate random-parameter negative binomial model under either dependence structure, with a random-coefficient data simulator, diagnostics, and standard model methods. Two estimation engines are provided: an OpenMP/Rcpp maximum-simulated-likelihood engine and a TMB automatic-differentiation engine that additionally offers a Laplace approximation, with a common front end.

License MIT + file LICENSE

Encoding UTF-8

Depends R (>= 4.1)

Imports stats,
compiler,
maxLik,
numDeriv,
randtoolbox,
MASS,
pbivnorm,
Rcpp,
parallel,
graphics,
TMB

LinkingTo Rcpp,
TMB,
RcppEigen

Suggests testthat (>= 3.1.5),
knitr,

```

rmarkdown,
pkgload
SystemRequirements GNU make
Config/testthat/edition 3
VignetteBuilder knitr
Roxygen list(markdown = TRUE)
LazyData false
Config/roxygen2/version 8.0.0

```

Contents

bnb_boundary_tests	3
bnb_elasticities	4
bnb_gof	5
bnb_marginal_effects	5
bnb_residual_checks	7
copula	7
fit_bnb	8
fit_rpbnb	10
fit_rpbnb_tmb	12
lr_test	16
plot.bnb_fit	18
plot.rpbnb_fit	19
predict.rpbnb_tmb_fit	19
residuals.bnb_fit	20
residuals.rpbnb_fit	21
rpbnb	22
rpbnb_boundary_tests	28
rpbnb_build_info	30
rpbnb_control	31
rpbnb_elasticities	37
rpbnb_marginal_effects	38
rpbnb_threads	40
rpbnb_tmb_boundary_tests	40
rpbnb_tmb_control	44
rpbnb_tmb_dependence_profile	47
rpbnb_tmb_elasticities	48
rpbnb_tmb_marginal_effects	49
rpbnb_tmb_max_workload	51
simulate_bnb	52
simulate_rpbnb	53
simulate_rpbnb_copula	54
simulate_rpbnb_tmb	55
summary.rpbnb_tmb_fit	56

bnb_boundary_tests *Boundary-corrected LR tests for the NB2 dispersions of a fit_bnb() fit*

Description

Runs a likelihood-ratio test for each free NB2 dispersion (m_1 , m_2) of a fitted fixed-coefficient bi-variate NB model, against a restricted refit at that margin's exact Poisson limit ($m = 0$). $m = 0$ sits on the boundary of the parameter space, so the natural-scale summary reports no Wald z/p for m_1/m_2 (see `fit_bnb()`); the valid test is `lr_test()`'s boundary = TRUE 50:50 chi-square mixture. This is the fixed-coefficient counterpart of `rpbnb_boundary_tests()` – the same test, restricted to the one boundary-null group `fit_bnb()` has (no random-coefficient scales).

Usage

```
bnb_boundary_tests(fit, data, control = rpbnb_control(compute_se = FALSE))
```

Arguments

fit	A converged bnb_fit (from <code>fit_bnb()</code>) with dependence = "independence" or "famoye". Not supported for a <code>copula()</code> dependence: <code>fit_bnb()</code> itself refuses to combine poisson_1/poisson_2 with a copula dependence, so there is no Poisson-limit restriction to test.
data	The data frame the model was fit on. Required – the fit object does not store it, and every restricted model is refit on it.
control	An <code>rpbnb_control()</code> for the restricted refits. Defaults to <code>compute_se = FALSE</code> (the LR test needs only logLik and the degrees of freedom).

Details

Each restricted refit is warm-started from the full fit's own coefficients (`start = fit$coef`); `fit_bnb()` pins the tested margin's `log_m` to the Poisson placeholder regardless of what `start` supplies for it, so passing the full vector is safe and needs no per-test trimming. Unlike `rpbnb_boundary_tests()`, there is no warm-start "polish" step for a negative LR statistic: `fit_bnb()`'s likelihood is exact (glm.nb for independence, closed-form-gradient BFGS for Famoye), not simulated, so a warm-started restricted refit does not out-climb the full fit's own optimum the way a simulated objective occasionally can.

Value

An object of class `bnb_boundary_tests` (a data frame with columns Parameter, LR, df, p.value, Signif, one row per tested margin) and a print method. A margin already fit at poisson_1/poisson_2 = TRUE has no free dispersion left to test and is skipped, matching `rpbnb_boundary_tests()`'s dispersion-test behaviour.

See Also

`lr_test()`, `fit_bnb()`, `rpbnb_boundary_tests()`

Examples

```
d <- read.csv(system.file("extdata", "rwm1984_clean.csv", package = "rpbnb"))
fit <- fit_bnb(docvis ~ outwork, hospvis ~ outwork, data = d,
               dependence = "famoye")
bnb_boundary_tests(fit, d)
```

bnb_elasticities	<i>Elasticities and semi-elasticities for a bivariate NB model</i>
------------------	--

Description

For continuous regressors, the elasticity is $\beta_j E[x_j]$; for binary (0/1) regressors, the semi-elasticity is $\exp(\beta_j) - 1$ (proportional change moving from 0 to 1).

Usage

```
bnb_elasticities(
  fit,
  which = c("y1", "y2", "both"),
  type = c("AME", "MEM"),
  vars = NULL,
  include_intercept = FALSE,
  digits = 4,
  print_output = TRUE
)
```

Arguments

fit	A bnb_fit object from <code>fit_bnb()</code> .
which	Which margin(s): "y1", "y2", or "both".
type	"AME" (uses sample mean of x) or "MEM" (uses mean design row).
vars	Optional variable names to restrict output.
include_intercept	Logical; include the intercept term.
digits	Number of decimal places for printed output.
print_output	Logical; if FALSE, suppress printing.

Value

A data frame (single margin) or a named list of data frames (both), each with columns Name, Estimate, StdErr, z, p, Signif, var_type.

Examples

```
d <- read.csv(system.file("extdata", "rwm1984_clean.csv", package = "rpbnb"))
fit <- fit_bnb(docvis ~ outwork + age, hospvis ~ outwork, data = d,
               dependence = "famoye")
bnb_elasticities(fit, which = "both", type = "AME")
```

bnb_gof

*Goodness of fit for a bivariate NB model***Description**

Reports log-likelihood, AIC, BIC, and four pseudo-R-squared measures (McFadden, McFadden adjusted, Cox-Snell, Nagelkerke) relative to an intercept-only null model refit with the same dependence structure (copula fits rebuild the copula from `fit$cop_family`). A margin fitted with an `offset()` keeps that offset in the null (an intercept-plus-offset model), so the null shares the full model's exposure/sampling structure and the pseudo-R-squared values remain comparable. Pseudo-R-squared values are returned raw and are not clamped to $[0, 1]$; a negative value flags a full model that fits worse than the null.

Usage

```
bnb_gof(fit, digits = 4, print_output = TRUE)
```

Arguments

`fit` A `bnb_fit` object from `fit_bnb()`.
`digits` Number of decimal places for printed output.
`print_output` Logical; if FALSE, suppress printing and return the result invisibly.

Value

Invisibly, a list with `n`, `k`, `logLik_full`, `logLik_null`, AIC, BIC, a named numeric vector `pseudoR2`, and the `null_fit`.

Examples

```
d <- read.csv(system.file("extdata", "rwm1984_clean.csv", package = "rpbnb"))
fit <- fit_bnb(docvis ~ outwork + age, hospvis ~ outwork, data = d,
               dependence = "famoye")
bnb_gof(fit)
```

bnb_marginal_effects

*Marginal effects for a bivariate NB model***Description**

Average marginal effects (AME) or marginal effects at the mean (MEM) for each margin. Continuous and binary (0/1) regressors are auto-detected; continuous effects use $\partial E[Y]/\partial x_j = \beta_j \mu$ and binary effects use $E[Y|x_j = 1] - E[Y|x_j = 0]$.

Usage

```
bnb_marginal_effects(
  fit,
  which = c("y1", "y2", "both", "all"),
  type = c("AME", "MEM"),
  vars = NULL,
  include_intercept = FALSE,
  digits = 4,
  print_output = TRUE
)
```

Arguments

<code>fit</code>	A <code>bnb_fit</code> object from <code>fit_bnb()</code> .
<code>which</code>	Which margin(s): "y1", "y2", "both", or "all".
<code>type</code>	"AME" (average marginal effect) or "MEM" (effect at the mean).
<code>vars</code>	Optional variable names or indices to restrict output.
<code>include_intercept</code>	Logical; include the intercept term.
<code>digits</code>	Number of decimal places for printed output.
<code>print_output</code>	Logical; if FALSE, suppress printing.

Details

For a model fitted with an `offset()`, absolute marginal effects scale with the fitted mean $\mu = \exp(x'\beta + \text{offset})$: AME uses each observation's training offset, and MEM uses the mean training offset.

Value

A data frame (single margin) or a named list of data frames (both), each with columns Name, Estimate, StdErr, z, p, Signif, var_type.

Examples

```
d <- read.csv(system.file("extdata", "rwm1984_clean.csv", package = "rpbnb"))
fit <- fit_bnb(docvis ~ outwork + age, hospvis ~ outwork, data = d,
  dependence = "famoye")
bnb_marginal_effects(fit, which = "y1", type = "AME")
```

bnb_residual_checks	<i>Residual checks for a bivariate NB model</i>
---------------------	---

Description

Formal residual diagnostics for a `fit_bnb()` or `fit_rpbnb()` fit: a normality test on the randomized quantile residuals (Shapiro-Wilk for $n \leq 5000$, else Kolmogorov-Smirnov vs $N(0,1)$); the NB2 dispersion statistic ($\text{sum}(\text{pearson}^2) / (n - k)$) per margin; the cross-margin RQR correlation (a check that the Famoye/copula dependence captured the association); an outlier count/index list ($|\text{RQR}| > \text{outlier_z}$); and a composite misspecification verdict combining these signals.

Usage

```
bnb_residual_checks(
  fit,
  seed = NULL,
  outlier_z = 3,
  digits = 4,
  print_output = TRUE
)
```

Arguments

<code>fit</code>	A <code>bnb_fit</code> or <code>rpbnb_fit</code> object.
<code>seed</code>	Optional integer seed for the RQR randomization.
<code>outlier_z</code>	Absolute-RQR threshold flagging an outlier (default 3).
<code>digits</code>	Decimal places for printed output.
<code>print_output</code>	Logical; if FALSE, suppress printing.

Value

Invisibly, an object of class `bnb_residual_checks`.

copula	<i>Specify a copula dependence structure</i>
--------	--

Description

Pass the result as the dependence argument to `fit_bnb()` or `fit_rpbnb()` to estimate a discrete-copula bivariate NB model instead of the default Famoye/Sarmanov dependence, or as the copula argument to `simulate_rpbnb_copula()` to simulate from one. The joint pmf is built from the two NB2 marginal CDFs and the chosen copula CDF (rectangle differencing); see the package vignette for the "Copula dependence" section.

Usage

```
copula(family = c("frank", "normal", "kimeldorf"), par = NULL)
```

Arguments

family One of "frank", "normal" (Gaussian), or "kimeldorf" (Clayton).

par Optional native dependence parameter (Frank's theta, the Gaussian rho, or Clayton's theta) used by `simulate_rpbnb_copula()` to generate data. Ignored by `fit_bnb()` and `fit_rpbnb()`, which estimate the parameter from the data.

Value

An object of class `rpbnb_copula`.

Examples

```
copula("frank")
copula("normal", par = 0.3)
copula("kimeldorf")
```

fit_bnb

Fit a bivariate negative binomial regression model

Description

Fit a bivariate negative binomial regression model

Usage

```
fit_bnb(
  formula_1,
  formula_2,
  data,
  dependence = c("independence", "famoye"),
  start = NULL,
  control = rpbnb_control(),
  poisson_1 = FALSE,
  poisson_2 = FALSE,
  boundary_tests = FALSE
)
```

Arguments

formula_1, formula_2	Model formulas for the two count outcomes. An equation-specific <code>offset()</code> term (e.g. <code>y ~ x + offset(log(exposure))</code>) is supported on every dependence path: the offset enters that margin's linear predictor additively (<code>mean exp(x'beta + offset)</code>) during estimation, and is carried through the stored fitted means and both <code>predict()</code> methods.
data	A data frame.
dependence	Dependence structure: "independence" (two univariate NB2 margins), "famoye" (Famoye/Sarmanov bivariate NB), or a <code>copula()</code> object (Frank / Gaussian / Clayton discrete-copula bivariate NB; the dependence parameter is estimated).
start	Optional starting parameter vector. May be positional (length equal to the number of parameters) or named; a named vector is reordered to the canonical parameter order and a named partial vector is merged into the defaults (unknown or duplicate names are rejected). When <code>start</code> is <code>NULL</code> , the famoye path uses a multi-start policy: it optimizes from both an all-zero mean-coefficient start and marginal <code>glm.nb</code> starts and keeps the better converged objective (the frozen-bounds gradient makes the objective start-sensitive and neither start dominates).
control	An <code>rpbnb_control()</code> object. The famoye and copula estimators both use BFGS, the only optimizer <code>control\$method</code> accepts. One control object serves every estimator in the package; settings this one does not read – <code>se_method</code> , <code>n_cores</code> , <code>halton_burn</code> , <code>draws_hessian</code> , and the TMB knobs – are ignored and listed by <code>print()/summary()</code> of the fit.
poisson_1, poisson_2	Fit the corresponding margin at its exact Poisson limit (NB2 dispersion $m = 0$) instead of estimating the dispersion. The margin's <code>log_m</code> is held fixed, so it is not a free parameter and the fit is a properly nested restriction of the NB model – pair it with <code>lr_test()</code> (<code>boundary = TRUE</code>) to test for overdispersion. This is the exact $m = 0$ restriction at any fitted mean: the margin's log-pmf is <code>dpois</code> and its Famoye dependence constant is <code>exp(-d*mu)</code> (the $m \rightarrow 0$ limit), not an NB2 at a tiny pinned dispersion. The famoye and independence paths are both exact (the independence path fits a Poisson GLM margin). Not supported with a <code>copula()</code> dependence.
boundary_tests	Run <code>bnb_boundary_tests()</code> on the converged fit and attach the result as <code>fit\$boundary_tests</code> , so <code>print()/summary()</code> fill in the <code>m1/m2 LR/df/p</code> columns instead of leaving them blank (see <code>bnb_boundary_tests()</code> for what the test is and why a Wald <code>z/p</code> is not valid there). Default <code>FALSE</code> : this costs up to two extra refits (one per free margin), so it is opt-in rather than automatic. Not supported with a <code>copula()</code> dependence, for the same reason <code>poisson_1/poisson_2</code> are not.

Value

An object of class `bnb_fit`.

Examples

```
d <- read.csv(system.file("extdata", "rwm1984_clean.csv", package = "rpbnb"))
```

```

fit <- fit_bnb(docvis ~ outwork, hospvis ~ outwork, data = d,
              dependence = "famoye")
summary(fit)

# Overdispersion test for margin 1 (H0: m1 = 0, Poisson)
fit_p1 <- fit_bnb(docvis ~ outwork, hospvis ~ outwork, data = d,
                  dependence = "famoye", poisson_1 = TRUE)
lr_test(fit_p1, fit, boundary = TRUE)

# Same test for both margins, run automatically and attached to the fit
fit2 <- fit_bnb(docvis ~ outwork, hospvis ~ outwork, data = d,
                dependence = "famoye", boundary_tests = TRUE)
fit2$boundary_tests

# Gaussian copula dependence instead of Famoye/Sarmanov
fit_cop <- fit_bnb(docvis ~ outwork, hospvis ~ outwork, data = d,
                  dependence = copula("normal"))
fit_cop$cop_tau # estimated Kendall's tau

```

fit_rpbnb

Fit a bivariate random-parameter negative binomial model

Description

Maximum simulated likelihood estimation with normal random coefficients and Famoye/Sarmanov dependence. Random coefficients are selected per equation by name via `random_1` / `random_2`.

Usage

```

fit_rpbnb(
  formula_1,
  formula_2,
  data,
  random_1 = NULL,
  random_2 = NULL,
  draws = 400,
  draw_type = "halton",
  seed = 1234,
  start = NULL,
  control = rpbnb_control(),
  dependence = "famoye",
  poisson_1 = FALSE,
  poisson_2 = FALSE,
  .fixed = NULL,
  .opt_draws = NULL
)

```

Arguments

formula_1, formula_2	Model formulas for the two count outcomes. An equation-specific <code>offset()</code> term (e.g. <code>y ~ x + offset(log(exposure))</code>) is supported: the offset enters that margin's linear predictor additively (integrated mean $E[\exp(x'\beta + \text{offset})]$) during estimation and is carried through the stored fitted means and <code>predict()</code> .
data	A data frame.
random_1, random_2	Random coefficients per equation. Either a character vector of <code>model.matrix</code> column names (all Normal), or a named list whose values are a distribution name ("normal", "lognormal", "uniform", "triangular") or a list <code>list(dist = ..., sign = ...)</code> (sign is -1/1 and lognormal-only). NULL means all-fixed for that equation.
draws	Number of simulation draws for the optimization.
draw_type	Quasi-random draw type. Only "halton" is supported in this version.
seed	Random seed for the simulation draws.
start	Optional starting parameter vector.
control	An <code>rpbnb_control()</code> object. Estimation uses BFGS, the only optimizer <code>control\$method</code> accepts. One control object serves every estimator in the package; settings this one does not read – the TMB knobs (<code>gradtol</code> , <code>restarts</code> , <code>max_threads</code> , <code>max_workload</code> , <code>parallel_tape</code>), <code>hessian</code> , and <code>draws_hessian</code> – are ignored and listed by <code>print()/summary()</code> of the fit rather than rejected.
dependence	Dependence structure: "famoye" (default; Famoye/Sarmanov) or an <code>copula()</code> object for copula dependence (Frank / Gaussian / Clayton). Both paths use the multithreaded (OpenMP) C++ simulated likelihood; the copula path is more numerically expensive per evaluation (discrete-copula pmf + per-draw NB CDF corners), so fits typically take noticeably longer than the Famoye path at comparable draws/n. Random coefficients on 0/1 dummy regressors are weakly identified under the copula path (NB dispersion trades off against the random-coefficient scale); prefer random coefficients on continuous regressors.
poisson_1, poisson_2	Fit the corresponding margin at its exact Poisson limit (NB2 dispersion $m = 0$): the margin's <code>log_m</code> is held fixed, so it is not a free parameter and the fit is a properly nested restriction of the NB model. Pair with <code>lr_test()</code> (boundary = TRUE) to test a margin for overdispersion ($H_0: m = 0$). The simulated likelihood uses the exact $m = 0$ branch – the per-draw margin log-pmf is <code>dpois</code> and its Famoye dependence constant is $\exp(-d \cdot \mu)$ – so it is accurate at any fitted mean, not a fixed-dispersion approximation. Supported with both Famoye/Sarmanov and <code>copula()</code> dependence (the copula path uses the same exact $m = 0$ branch).
.fixed	Internal. A named numeric vector of parameters (in the optimization/log-scale parameterization) to pin at the supplied values and hold fixed during estimation. Used by <code>rpbnb_boundary_tests()</code> to construct scale-zero (SD boundary) restricted fits; not intended for direct use.
.opt_draws	Internal. A list <code>list(Z1, Z2)</code> of uniform Halton draw matrices to use verbatim instead of generating them, so a restricted refit reuses a full fit's draws (common

random numbers). Used by `rpbnb_boundary_tests()`; not intended for direct use.

Value

An object of class `rpbnb_fit`. Under Famoye dependence three fields describe the admissible lambda interval: `bounds` is the interval the optimized likelihood actually used, frozen at the starting values (its width also feeds `summary()`'s delta-method standard error for lambda); `bounds_at_optimum` is the interval admissible at the fitted parameters, recomputed after optimization for validation only; and `lambda_admissible` records whether lambda — the fitted value, mapped through `bounds` — lies inside `bounds_at_optimum`. FALSE means the optimizer escaped the valid region (a warning is raised) and the fit should be re-run from starting values closer to the optimum. For normal or lognormal random coefficients loaded in both margins the bound is the constant `c(-1, 1)`, the two intervals coincide, and escape cannot occur.

Examples

```
sim <- simulate_rpbnb(n = 600,
  beta1 = c("(Intercept)" = 0.2, x1 = 0.4),
  beta2 = c("(Intercept)" = 0.1, x1 = -0.3),
  random_1 = list(x1 = list(sd = 0.5)),
  dispersion = c(m1 = 0.4, m2 = 0.5), seed = 1)
fit <- fit_rpbnb(y1 ~ x1, y2 ~ x1, data = sim$data, random_1 = "x1",
  draws = 100, control = rpbnb_control(compute_se = FALSE))
coef(fit)

# Copula dependence instead of Famoye/Sarmanov (slower; fewer draws here
# for a quick example -- use more in practice)
sim_cop <- simulate_rpbnb_copula(n = 600,
  beta1 = c("(Intercept)" = 0.2, x1 = 0.4),
  beta2 = c("(Intercept)" = 0.1, x1 = -0.3),
  random_1 = list(x1 = list(sd = 0.3)),
  dispersion = c(m1 = 0.4, m2 = 0.5),
  copula = copula("normal", par = 0.5), seed = 1)
fit_cop <- fit_rpbnb(y1 ~ x1, y2 ~ x1, data = sim_cop$data, random_1 = "x1",
  dependence = copula("normal"), draws = 100,
  control = rpbnb_control(compute_se = FALSE))
tanh(coef(fit_cop)[["z_theta"]]) # estimated copula rho
```

fit_rpbnb_tmb

Fit a bivariate random-parameter negative binomial model (TMB)

Description

Maximum simulated likelihood via TMB (Template Model Builder) with automatic differentiation. Supports Famoye/Sarmanov and discrete-copula (Frank, Gaussian, Clayton) dependence.

Usage

```
fit_rpbmb_tmb(
  formula_1,
  formula_2,
  data,
  random_1 = NULL,
  random_2 = NULL,
  draws = 400L,
  seed = 1234L,
  start = NULL,
  dependence = "famoye",
  control = rpbmb_tmb_control(),
  inference = c("full", "diag", "none"),
  keep = c("postfit", "compact", "full"),
  poisson_1 = FALSE,
  poisson_2 = FALSE,
  method = c("sml", "laplace"),
  disable_parallel_gaussian = FALSE,
  force_parallel_gaussian = NULL,
  .fixed = NULL
)
```

Arguments

formula_1, formula_2	Model formulas for the two count outcomes.
data	A data frame.
random_1, random_2	Random coefficient specs per equation. NULL (all fixed), character vector (all Normal), or named list with per-variable distribution specs ("normal", "lognormal", "uniform", "triangular").
draws	Number of Halton simulation draws. Under method = "sml" this sets the simulation grid the likelihood is averaged over. Tape size scales with nrow(data) * draws unless the fit draw-chunks (see control\$tape_chunks at rpbmb_control()): when the weighted workload exceeds control\$max_workload, the fit is split into several equal-sized draw chunks replayed over one smaller TMB tape, dropping peak memory to nrow(data) * ceiling(draws / chunks) at the cost of somewhat slower gradient evaluations – exact for the requested draws, not an approximation. Under method = "laplace" draws does not affect the likelihood, the tape, or chunking; its only remaining job is sizing the Halton grid used for the post-estimation averaging in predict() and the marginal-effect functions. (The frozen Famoye admissible-lambda interval is derived from the random coefficients' support, independent of draws under either estimator.)
seed	Random seed for draws.
start	Optional starting parameter vector (named or positional).
dependence	How the two margins are linked. "famoye" (default) is Famoye/Sarmanov dependence: a single bounded association parameter (lam), admissible inter-

	val frozen at the starting values. A <code>copula()</code> object joins the margins with a discrete copula instead – <code>copula("frank")</code>, <code>copula("normal")</code> (Gaussian), or <code>copula("kimeldorf")</code> (Clayton) – each with its own native dependence parameter, always estimated (<code>copula()</code>'s <code>par</code> argument is for the simulators only). Copula evaluation costs more per iteration than Famoye at comparable draws/n, and the dependence family also changes peak memory under this engine – see <code>max_workload</code> at <code>rpbmb_control()</code> for the measured per-family weights. "independence" – two separate NB2 margins, no association parameter at all – is only available here; the classic engine (<code>fit_rpbmb()</code>) has no independence path for random-parameter models (use <code>fit_bmb()</code> for a fixed-coefficient independence model). Whichever dependence is chosen, a random coefficient on a 0/1 dummy regressor present in both equations is weakly identified (NB dispersion trades off against the random-coefficient scale); prefer a continuous regressor for a shared random coefficient when one is available.
control	An <code>rpbmb_control()</code> object (<code>rpbmb_tmb_control()</code> is a retained alias that returns the same object). One control object serves every estimator in the package; settings this engine does not read – <code>se_method</code>, <code>hessian</code>, <code>compute_se</code>, <code>method</code>, <code>hess_eps</code>, <code>hess_r</code>, <code>draws_hessian</code> – are ignored and listed by <code>print()/summary()</code> of the fit rather than rejected.
inference	Inference storage: "full" for a full covariance, "diag" for standard errors only, or "none" to skip Hessian calculations. In diagonal mode, <code>vcov()</code> returns NA for covariance cross-terms.
keep	Retained fit state: "postfit" keeps data needed for marginal effects, "compact" keeps estimates only, and "full" also retains the TMB objective and responses. Low-level diagnostics that access <code>fit\$obj</code> require "full".
poisson_1, poisson_2	Fit the corresponding margin at its exact Poisson limit (NB2 dispersion $m = 0$, held fixed). Available for every dependence structure, copulas included. This is the remedy for a dispersion that collapses toward zero on its own: at $m \sim 1e-7$ the NB2 size $1/m$ is $\sim 1e7$ and the marginal CDF is computed as a difference of logarithms agreeing to eleven digits, which leaves the value intact but the curvature – and so the standard errors – numerically worthless. Pinning the limit removes the parameter instead of estimating it into that regime.
method	Estimator for the random-coefficient integral. "sml" (default) uses simulated maximum likelihood over Halton draws. "laplace" uses TMB's Laplace approximation with a sparse Hessian over one latent vector per observation, which removes draws from the memory cost and makes tape size scale with <code>nrow(data)</code> alone. The two are different approximations to the same integral. They agree asymptotically but need not agree closely on any given dataset, so a Laplace fit is not a drop-in reproduction of an SML fit. "laplace" supports "normal" and "lognormal" random coefficients only, and requires at least one random coefficient. <code>summary()</code> reports AIC and BIC for either estimator, but under "laplace" they are computed from an approximated marginal likelihood rather than the exact one, so an AIC from a Laplace fit is not meaningful to compare against an AIC from an SML fit of the same model.

<code>disable_parallel_gaussian</code>	Opt-out that restricts a Gaussian-copula (<code>dependence = copula("normal")</code>) fit to one thread. Default FALSE: the requested <code>control\$n_cores/control\$parallel_tape</code> are honored for every dependence family, Gaussian included. Multithreaded Gaussian evaluation used to crash the R process (SIGSEGV) whenever a fit realized more threads than an earlier fit in the same session – TMB’s REGISTER_ATOMIC cache sizes its per-thread array once, at first initialization – and Gaussian fits were therefore capped to one thread by default. That defect was fixed in 0.4.4 (<code>src/rpbmb_tmb.cpp</code> sizes the atomic for the machine’s full processor count at initialization; verified at 2/4/8/16 threads), and since 0.4.6 parallel evaluation is the default. Setting <code>disable_parallel_gaussian = TRUE</code> restores the old single-thread behaviour – a kill-switch for an unusual TMB/OpenMP build where the registered atomic still misbehaves. Ignored (threads never restricted) for every other dependence structure.
<code>force_parallel_gaussian</code>	Deprecated (pre-0.4.6 polarity); use <code>disable_parallel_gaussian</code> instead. Multithreaded Gaussian evaluation is now the default, so TRUE (the old opt-in) is a no-op beyond its deprecation warning, and an explicit FALSE – which used to select the single-thread cap – is honored as <code>disable_parallel_gaussian = TRUE</code> . Any non-NULL value warns.
<code>.fixed</code>	Internal. A named numeric vector of parameters (in the optimization parameterization, e.g. <code>c("log_sd1:x" = -20)</code>) to pin at the supplied values and hold fixed during estimation, so they leave the free-parameter count and <code>logLik()</code> ’s df. Used by <code>rpbmb_tmb_boundary_tests()</code> to construct scale-zero restricted fits; not intended for direct use. The classic engine’s <code>fit_rpbmb()</code> carries an identically named argument.

Value

An object of class `rpbmb_tmb_fit`. The `sdreport` field is a compact package-owned summary and does not retain a second TMB tape.

`optimizer` is the `nlminb()` return value, with three fields added: `restarts` counts how many times the solve was restarted from its own answer to clear `control$gradtol`, `max_abs_gradient` is the score norm actually achieved, and `gradient_tolerance` is the threshold it was judged against. `convergence` and `message` come from `nlminb`’s relative *function* test and can report success at a point that is not stationary, so `max_abs_gradient` is the field to check before trusting a standard error.

`method` records which estimator produced the fit, echoing the method argument (`"sml"` or `"laplace"`). Both `print()` and `summary()` display it, so two fits of the same model under different estimators are distinguishable in their printed output.

`boundary_report` is a character vector naming the reported quantities whose estimates are pinned against an implementation bound – the frozen Famoye interval, the Frank overflow guard, or an `exp()` clamp – or whose delta-method derivative has collapsed numerically. Those estimates are set by the implementation rather than identified by the data, so their standard errors and covariances are reported as NA and a warning is raised. An empty vector means no such constraint bound.

`boundary_sides` names, for each entry of `boundary_report`, which end was reached: `"lower"`, `"upper"`, or `"degenerate"` when the derivative collapsed rather than a bound being hit. The two

sides call for opposite remedies – a lower dispersion clamp means the margin is effectively Poisson, an upper one means it is degenerately over-dispersed – so this is retained on the fit rather than only mentioned in the warning.

`lambda_bounds` is a named numeric vector `c(lower =, upper =)` giving the admissible Famoye dependence interval, and `NULL` for every other dependence structure. The bounds *used by the likelihood* are computed once at the starting values and held fixed for the whole fit. A `lam` estimate at either end is therefore an artefact of the starting values rather than a property of the data, which is why the field is exposed. `rpbnb_tmb_dependence_profile` reports a likelihood-based interval where the delta-method standard error collapses to `NA`, but for Famoye that interval is mapped through this same frozen box and so cannot escape it: widen the box by refitting from better starting values before treating such an interval as informative.

`lambda_bounds_at_optimum` is the interval admissible at the *fitted* parameters, recomputed after optimization for validation only — it never enters the likelihood. `lambda_admissible` records whether the fitted `lam` (mapped through the frozen `lambda_bounds`, i.e. the value the objective actually used) lies inside it; `FALSE` means the optimizer left the region where the joint pmf is a valid probability model, a warning was raised, and the fit should not be interpreted — refit from starting values closer to the optimum. When the bound is parameter-independent (normal or lognormal coefficients loaded in both margins, where it is the constant `c(-1, 1)`), the two intervals coincide and the flag is trivially `TRUE`. Both fields are `NULL/NA` outside Famoye dependence.

Examples

```
## Not run:
sim <- simulate_rpbnb_tmb(n = 300,
  beta1 = c("(Intercept)" = 0.2, x1 = 0.4),
  beta2 = c("(Intercept)" = 0.1, x1 = -0.3),
  dispersion = c(m1 = 0.4, m2 = 0.5), seed = 1)
fit <- fit_rpbnb_tmb(y1 ~ x1, y2 ~ x1, data = sim$data, draws = 100)
coef(fit)

## End(Not run)
```

lr_test

Likelihood-ratio test between two nested model fits

Description

Compares a restricted fit against a full (nesting) fit by the likelihood-ratio statistic. Works for any `fit_bnb()` / `fit_rpbnb()` objects, since both carry a `logLik()` with a "df" attribute equal to the number of estimated parameters.

Usage

```
lr_test(restricted, full, boundary = FALSE)
```

Arguments

restricted	The smaller (restricted) fit – fewer estimated parameters.
full	The larger (full) fit that nests restricted.
boundary	Logical. When the restriction pins a variance/dispersion-type parameter (a random-coefficient SD, or NB2 dispersion m) to its zero boundary, the null distribution is not a plain chi-square. Set <code>boundary = TRUE</code> to use the 50:50 mixture of <code>chisq(df)</code> and <code>chisq(df - 1)</code> (Self & Liang, 1987); for a single boundary parameter ($df = 1$) this halves the naive p-value. Default <code>FALSE</code> (ordinary interior restriction, e.g. dropping a fixed covariate). The mixture is exact only for a single parameter on the boundary; simultaneous boundary restrictions of several parameters need different weights and are not handled.

Details

The restricted model is one you fit yourself with a term removed – for example dropping a name from `random_1` (testing a random-coefficient SD), or a plain NB / independence fit (testing an NB2 dispersion). This is the statistically appropriate replacement for the Wald z/p that the natural-scale summary suppresses on positive scale/dispersion parameters.

A likelihood-ratio test requires two *maximized* likelihoods. When either argument is a package fit (`bnb_fit` / `rpbnb_fit`) that records a failed optimization (`convergence$converged = FALSE`), `lr_test()` errors rather than returning a p-value from an unfinished fit – the sign of the statistic cannot establish convergence. Generic objects that only carry a `logLik()` (no convergence record) are still accepted, but their convergence cannot be validated and is the caller's responsibility.

Value

An object of class `rpbnb_lrtest` with the LR statistic, degrees of freedom `df`, p.value, the two log-likelihoods and their `df`, and the boundary flag. Has a `print` method.

Famoye caveat for bounded random-coefficient distributions

Each Famoye fit maximizes its likelihood over an admissible `lambda` interval frozen at that fit's own starting values. With normal or lognormal random coefficients loaded in both margins the interval is the constant `c(-1, 1)` for every fit, so the comparison is unaffected. With uniform or triangular coefficients (or a single varying margin) the bound moves with the parameters, so two fits frozen at different starting values maximize over slightly different `lambda` ranges and the LR statistic inherits that second-order discrepancy. Check `lambda_admissible` on both fits before relying on a comparison in that regime.

References

Self, S. G. and Liang, K.-Y. (1987). Asymptotic properties of maximum likelihood estimators and likelihood ratio tests under nonstandard conditions. *JASA* 82(398), 605–610.

Examples

```
sim <- simulate_rpbmb(n = 600,
  beta1 = c("(Intercept)" = 0.2, x1 = 0.4),
```

```

beta2 = c("(Intercept)" = 0.1, x1 = -0.3),
random_1 = list(x1 = list(sd = 0.5)),
dispersion = c(m1 = 0.4, m2 = 0.5), seed = 1)
ctrl <- rpbnb_control(compute_se = FALSE)
full <- fit_rpbnb(y1 ~ x1, y2 ~ x1, data = sim$data, random_1 = "x1",
                 draws = 100, control = ctrl)
rest <- fit_rpbnb(y1 ~ x1, y2 ~ x1, data = sim$data,
                 draws = 100, control = ctrl) # no random coefficient
lr_test(rest, full, boundary = TRUE)         # test sd(x1) = 0

```

plot.bnb_fit

Residual diagnostic plots for a bivariate NB model

Description

Four base-graphics panels per margin: residuals-vs-fitted, a normal QQ plot of the randomized quantile residuals, a histogram of the RQR with an $N(0,1)$ overlay, and a scale-location plot. The QQ and histogram panels always use RQR (only these are approximately $N(0,1)$ under a correct count model).

Usage

```

## S3 method for class 'bnb_fit'
plot(
  x,
  margin = c("both", "y1", "y2"),
  which = 1:4,
  resid_type = "quantile",
  seed = NULL,
  ...
)

```

Arguments

x	A bnb_fit object from fit_bnb() .
margin	Which margin to plot: "both" (default), "y1", or "y2".
which	Integer subset of panels 1:4 (1 = residuals-vs-fitted, 2 = QQ, 3 = histogram, 4 = scale-location).
resid_type	Residual type for panels 1 and 4: "quantile" (default), "pearson", "deviance", or "response".
seed	Optional integer seed for the RQR randomization.
...	Unused.

Value

NULL, invisibly (called for the side effect of drawing).

plot.rpbnb_fit	<i>Residual diagnostic plots for a random-parameter bivariate NB model</i>
----------------	--

Description

Four base-graphics panels per margin, as for `plot.bnb_fit()`, built on the mixture-based randomized quantile residuals. `resid_type = "deviance"` is not available for `rpbnb_fit`.

Usage

```
## S3 method for class 'rpbnb_fit'
plot(
  x,
  margin = c("both", "y1", "y2"),
  which = 1:4,
  resid_type = "quantile",
  seed = NULL,
  ...
)
```

Arguments

<code>x</code>	An <code>rpbnb_fit</code> object from <code>fit_rpbnb()</code> .
<code>margin</code>	Which margin to plot: "both" (default), "y1", or "y2".
<code>which</code>	Integer subset of panels 1:4.
<code>resid_type</code>	Residual type for panels 1 and 4: "quantile" (default), "pearson", or "response".
<code>seed</code>	Optional integer seed for the RQR randomization.
<code>...</code>	Unused.

Value

NULL, invisibly.

predict.rpbnb_tmb_fit	<i>Predict from a fitted bivariate count model</i>
-----------------------	--

Description

Predictions integrate over the retained simulation draws when random coefficients are present. Link predictions are the log of the integrated response mean, so `exp(predict(fit, type = "link"))` equals response predictions.

Usage

```
## S3 method for class 'rpbnb_tmb_fit'
predict(
  object,
  newdata = NULL,
  type = c("response", "link"),
  which = c("both", "y1", "y2"),
  ...
)
```

Arguments

object	A fitted rpbnb_tmb_fit object.
newdata	Optional data frame. If omitted, the retained fitting design is used; compact fits require newdata.
type	Either "response" or "link".
which	Return both margins or only "y1" or "y2".
...	Reserved for future use.

Value

A two-column numeric matrix for which = "both", otherwise a numeric vector.

residuals.bnb_fit	<i>Residuals for a bivariate NB model</i>
-------------------	---

Description

Per-margin residuals for a fixed-coefficient `fit_bnb()` model. Each margin is NB2 with fitted mean μ and dispersion m (size = $1/m$). "quantile" returns randomized quantile residuals (Dunn & Smyth 1996), which are approximately $N(0,1)$ under a correct model and are the recommended residual for normality-style diagnostics on count data.

Usage

```
## S3 method for class 'bnb_fit'
residuals(
  object,
  type = c("quantile", "pearson", "deviance", "response"),
  margin = c("both", "y1", "y2"),
  seed = NULL,
  ...
)
```

Arguments

object	A bnb_fit object from <code>fit_bnb()</code> .
type	Residual type: "quantile" (default), "pearson", "deviance", or "response".
margin	Which margin: "both" (default), "y1", or "y2".
seed	Optional integer seed for the quantile-residual randomization (ignored for other types); does not disturb the caller's RNG stream.
...	Unused.

Value

A numeric vector for a single margin, or a two-column data frame (y1, y2) for margin = "both".

residuals.rpbnb_fit *Residuals for a random-parameter bivariate NB model*

Description

Per-margin residuals for an `fit_rpbnb()` model. Each margin is a mixture of NB2 distributions over the random-coefficient draws. "quantile" returns randomized quantile residuals (Dunn & Smyth 1996) from the exact mixture predictive CDF (the recommended residual); "pearson" uses the exact mixture marginal variance. "deviance" is not defined for the mixture and errors.

Usage

```
## S3 method for class 'rpbnb_fit'
residuals(
  object,
  type = c("quantile", "pearson", "deviance", "response"),
  margin = c("both", "y1", "y2"),
  seed = NULL,
  ...
)
```

Arguments

object	An rpbnb_fit object from <code>fit_rpbnb()</code> .
type	Residual type: "quantile" (default), "pearson", or "response". "deviance" is not supported for rpbnb_fit.
margin	Which margin: "both" (default), "y1", or "y2".
seed	Optional integer seed for the quantile-residual randomization; does not disturb the caller's RNG stream.
...	Unused.

Value

A numeric vector for a single margin, or a two-column data frame (y1, y2) for margin = "both".

rpbnb

Fit a random-parameter bivariate NB model with either engine

Description

A common front end over the package's two estimation engines. `engine = "classic"` calls `fit_rpbnb()` (Rcpp/OpenMP simulated likelihood, `maxLik` BFGS); `engine = "tmb"` calls `fit_rpbnb_tmb()` (TMB automatic differentiation, `nlminb` with restart polish). Both fitters remain exported and can be called directly; with `standardize = FALSE` (the default) this wrapper adds nothing to the fit itself and returns exactly what the chosen fitter returns.

Usage

```
rpbnb(
  formula_1,
  formula_2,
  data,
  engine = c("classic", "tmb"),
  random_1 = NULL,
  random_2 = NULL,
  draws = 400,
  seed = 1234,
  start = NULL,
  dependence = "famoye",
  poisson_1 = FALSE,
  poisson_2 = FALSE,
  standardize = FALSE,
  continuous_vars = NULL,
  boundary_tests = FALSE,
  boundary_draws = NULL,
  control = NULL,
  ...
)
```

Arguments

<code>formula_1, formula_2</code>	Model formulas for the two count responses.
<code>data</code>	A data frame containing the model variables.
<code>engine</code>	Which fitter estimates the model. "classic" (default) calls <code>fit_rpbnb()</code> : a multithreaded (OpenMP) Rcpp simulated-likelihood engine optimized by <code>maxLik::maxLik(method = "BFGS")</code> with a numerical gradient. "tmb" calls <code>fit_rpbnb_tmb()</code> : it builds an automatic-differentiation tape with TMB (Template Model Builder) and optimizes it with <code>stats::nlminb</code> plus a restart polish – the exact gradient this gives tends to converge faster and more reliably than the classic engine's numerical one, especially under copula dependence. "tmb" is also the only engine

that accepts `dependence = "independence"`, the Laplace estimator (`method = "laplace"`, passed via `...` – see the "estimator" discussion below and `fit_rpbnb_tmb()`'s `method` argument), and memory-aware draw chunking for large draws; "classic" is the only one that accepts an `offset()` term in a formula. Both engines share one `rpbnb_control()` object and this function's cross-engine argument checking, but otherwise a fit is exactly what a direct call to the chosen fitter would return – see "Which arguments go with which engine" below for the full compatibility table, and the Return section for how the fit class depends on engine.

<code>random_1, random_2</code>	Random-coefficient specifications for each equation.
<code>draws</code>	Number of simulation draws.
<code>seed</code>	Random seed for the draw sequence.
<code>start</code>	Optional named or unnamed starting values.
<code>dependence</code>	How the two margins are linked. "famoye" (default) is Famoye/Sarmanov dependence: a single bounded association parameter (<code>lam</code>), with an admissible interval frozen at the starting values (see "Which arguments go with which engine" below). A <code>copula()</code> object joins the margins with a discrete copula instead – <code>copula("frank")</code> , <code>copula("normal")</code> (Gaussian), or <code>copula("kimeldorf")</code> (Clayton) – each with its own native dependence parameter, always estimated (<code>copula()</code> 's <code>par</code> argument is for the simulators only, not the fitters). Copula evaluation costs more per iteration than Famoye at comparable draws/n (discrete-copula pmf plus a per-draw NB CDF corner), and under <code>engine = "tmb"</code> the dependence family also changes peak memory – see <code>max_workload</code> at <code>rpbnb_control()</code> for the measured per-family weights. "independence" – two separate NB2 margins, no association parameter at all – is <code>engine = "tmb"</code> only; the fixed-coefficient <code>fit_bnb()</code> supports it under <code>engine = "classic"</code> , but the random-parameter classic engine (<code>fit_rpbnb()</code>) does not. Whichever dependence is chosen, a random coefficient on a 0/1 dummy regressor present in both equations is weakly identified (NB dispersion trades off against the random-coefficient scale); prefer a continuous regressor for a shared random coefficient when one is available.
<code>poisson_1, poisson_2</code>	Restrict the corresponding margin to its Poisson limit.
<code>standardize</code>	Centre and scale continuous predictors before fitting, and display fitted coefficients back-transformed to their original units (see "Automatic centring and scaling" above). Default FALSE.
<code>continuous_vars</code>	Optional character vector overriding automatic continuous-predictor detection when <code>standardize = TRUE</code> . Must be columns of data; ignored when <code>standardize = FALSE</code> .
<code>boundary_tests</code>	Which groups of parameters to LR-test after fitting, via <code>rpbnb_boundary_tests()</code> (<code>engine = "classic"</code>) or <code>rpbnb_tmb_boundary_tests()</code> (<code>engine = "tmb"</code>); the result is attached as <code>\$boundary_tests</code> and <code>summary()/print()</code> show the test (rather than NA, or rather than a Wald z) on the corresponding rows. Accepts: FALSE (default) or "none" – run nothing. Each restricted refit costs roughly another full fit.

- TRUE – c("sd", "dispersion"), the two boundary-null groups. This is what TRUE has always meant and it does not silently grow.
- a character vector, any of "sd" (random-coefficient scales), "dispersion" (the NB2 overdispersions m1, m2), "dependence" (the association parameter), or "all" for all three.

So `boundary_tests = "dispersion"` tests overdispersion only, `boundary_tests = c("dispersion", "dependence")` tests overdispersion and association without paying for one refit per random-coefficient scale, and `boundary_tests = "all"` tests everything. See "Boundary LR tests" below.

<code>boundary_draws</code>	Number of Halton simulation draws for the boundary tests' restricted refits, <code>engine = "tmb"</code> only. NULL (default) uses the main fit's draws. Ignored when no boundary test was requested (nothing to apply it to); an error if supplied under <code>engine = "classic"</code> alongside a boundary test (see "Boundary LR tests" above – that engine has no draws to override).
<code>control</code>	An <code>rpbnb_control()</code> object – one object for both engines, as of 0.4.1 (<code>rpbnb_tmb_control()</code> is a retained alias that returns the same thing). Settings the chosen engine does not read are ignored and listed by <code>print()/summary()</code> of the fit; <code>iterlim</code> and <code>print_level</code> , whose defaults differ between the engines, resolve to the chosen engine's own default unless you set them. Fields sharing a name are still not translated: <code>iterlim</code> is a <code>maxLik</code> BFGS limit under <code>engine = "classic"</code> and an <code>nlminb</code> limit under <code>"tmb"</code> , and <code>n_cores</code> is worker processes versus OpenMP threads.
<code>...</code>	Further arguments passed to the selected fitter. Names are validated against that fitter's formals; an argument belonging to the other engine, or an unrecognised name, is an error. Exception: the TMB tuning knobs method, <code>disable_parallel_gaussian</code>, and the deprecated <code>force_parallel_gaussian</code> are dropped with a warning (not an error) under <code>engine = "classic"</code>, so a call can switch engines without stripping them. The estimator, <code>method</code> (<code>engine = "tmb"</code> only, default <code>"sml"</code>), picks how the random-coefficient integral is approximated. <code>"sml"</code> is simulated maximum likelihood over draws Halton points – the same integral the classic engine approximates, but with an exact automatic-differentiation gradient. <code>"laplace"</code> instead uses TMB's Laplace approximation: a sparse-Hessian integration over one latent vector per observation that removes draws from the memory cost entirely (tape size then scales with <code>nrow(data)</code> alone, not <code>nrow(data) * draws</code>), at the cost of requiring at least one random coefficient and restricting it to <code>"normal"/"lognormal"</code> (<code>"uniform"</code> and <code>"triangular"</code> error under Laplace). The two estimators agree asymptotically but are different approximations to the same integral – not interchangeable point-for-point on a given dataset, and their <code>summary()</code> AIC/BIC are not comparable to each other for the same reason. See <code>fit_rpbnb_tmb()</code>'s <code>method</code> argument for the full detail, including how boundary tests and <code>predict()</code> behave differently under each.

Details

What it does add is argument checking across the two APIs. The engines do not take the same arguments, and passing one engine's argument to the other is a mistake that is easy to make and

expensive to notice — a standardized coefficient table printed under an "original units" heading looks perfectly plausible. Every extra argument is therefore matched by name against the selected fitter's own formal, and anything that does not belong is an error rather than a silently ignored ... entry.

Value

The engine-native fit object, identical to what a direct call to the underlying fitter would return: an object of class `rpbnb_fit` for engine = "classic", or `rpbnb_tmb_fit` for engine = "tmb". The class therefore depends on engine; test with `inherits(fit, "rpbnb_tmb_fit")` if you need to branch. No wrapper class is introduced, so every existing S3 method and post-estimation function works unchanged. With `standardize = TRUE`, two extra fields are attached — `$scaling` and `$continuous_vars` — and `print()/summary()` use them to display original-units coefficients. With any `boundary_tests` group requested, a third field `$boundary_tests` (the [rpbnb_boundary_tests\(\)](#) result) is attached, and `print()/summary()` use it to show the LR test on the corresponding rows. Both fitters also attach `$control_ignored` / `$control_engine`, the control settings the chosen engine did not read. Nothing else on the object changes.

Automatic centring and scaling

`standardize = TRUE` automates the pattern in `inst/dev/rpbnb_frank_open.R` and `inst/dev/tmb_rpbnb_frank_open.R`: continuous predictors are centred and scaled (mean 0, SD 1) before fitting, which keeps a bounded random-coefficient carrier from acting as a disguised random intercept (see those scripts' headers) and fixes the design matrix's conditioning when regressors span very different ranges. Continuous predictors are identified automatically — numeric, non-factor columns used by either formula with more than two distinct values, so 0/1 (or any two-level numeric) indicators are left alone — or supplied explicitly via `continuous_vars`. Variables that appear only inside an `offset()` are never standardized.

The fitted design itself (the stored $X1/X2$, $\mu1/\mu2$, simulation draws) stays on the standardized scale, exactly as in the two scripts above, so `predict()`, marginal effects, and boundary/LR tests keep working unchanged. Only the coefficient table `print()` and `summary()` display is affected: it is back-transformed to the covariates' original units by the exact affine chain rule (no refit, no numerical differentiation) and is the *only* coefficient table shown — there is no separate standardized-scale table to reconcile. `coef()/vcov()` still return the standardized-scale values that match the stored design; call `.rpbnb_orig_units()` (internal, mirrors what `print()` displays) if the numeric original-units vector is needed directly. The scaling actually used is stored on the fit as `$scaling` (a named list of `c(center=, scale=)`) and `$continuous_vars`.

Boundary LR tests

`boundary_tests` runs a likelihood-ratio test on the fit as soon as it converges (against the same data the fit itself used — the standardized copy, when `standardize = TRUE`) and attaches the result as `$boundary_tests`. It is a switch over three independent groups, so the cost is paid only for the parameters actually in question:

<code>boundary_tests</code>	tests	restricted refits
<code>FALSE / "none"</code>	nothing	0
<code>TRUE</code>	<code>c("sd", "dispersion")</code>	one per scale + one per free m
<code>"sd"</code>	random-coefficient scales	one per scale

"dispersion"	overdispersions m1, m2	one per free m
"dependence"	the association parameter	1
"all"	all three groups	all of the above

The random-coefficient SDs and the NB2 dispersions (m_1 , m_2) have a null that sits on the boundary of the parameter space, so an ordinary Wald z/p does not test it; the boundary test refits each restricted model instead and `summary()/print()` show that test's LR/df/p for those rows in the natural-scale block, in place of the NA they would otherwise carry.

The **dependence** test is the odd one out and is therefore opt-in even under `boundary_tests = TRUE`. Its null is "no association", i.e. the independence model, which for Famoye (λ), Frank (θ), and the Gaussian copula (ρ) sits in the *interior* of the parameter space – the Wald z those rows already show is valid, and the LR statistic is an ordinary chi-square(1), not the 50:50 boundary mixture. Only Clayton / Kimeldorf ($\theta > 0$) has a genuine boundary null there. Requesting it replaces the dependence row's Wald z/p with the LR test on both engines, since the two answer the same question.

Both engines test the same parameters via `rpbnb_boundary_tests()` (`engine = "classic"`) or `rpbnb_tmb_boundary_tests()` (`engine = "tmb"`). They differ only in how each restricted fit is constructed while preserving common random numbers: for a scale, the classic engine zeroes that coefficient's draw column while the TMB engine pins its `log_sd` and maps it out of the free parameters; for the dependence parameter, the TMB engine refits with its own `dependence = "independence"` family while the classic engine (which has no such fitter) pins the working-scale dependence parameter at its family's independence value.

A `message()` reports how many restricted refits are about to run before they start, unless `control$print_level` is 0; suppress it with `suppressMessages()` if needed.

Under `engine = "tmb"` with `method = "laplace"`, a restricted refit that fails to converge is automatically retried with both sides of that one LR estimated by `method = "sml"` rather than reported NA – some restrictions leave Laplace no valid optimum at all (see `rpbnb_tmb_boundary_tests()`'s `sml_fallback` argument, which is where to turn this off).

`disable_parallel_gaussian` (`engine = "tmb"` only, passed via ...) is forwarded to every restricted refit, so a Gaussian-copula fit that opted out of multithreading gets single-threaded boundary refits too – see `rpbnb_tmb_boundary_tests()`'s own `disable_parallel_gaussian` argument for why this needs forwarding at all (the fit object does not record whether the opt-out was used). The deprecated pre-0.4.6 `force_parallel_gaussian` is likewise accepted and mapped (see `fit_rpbnb_tmb()`).

Each restricted refit costs roughly as much as the original fit (more for a `copula()` dependence than for "famoye"; see `rpbnb_boundary_tests()`'s timing note), so this defaults to FALSE. `boundary_draws` (`engine = "tmb"` only) sets a draws for those refits other than the main fit's – e.g. more draws for a more precise boundary test without re-fitting the whole model at that draws. It has no classic-engine counterpart: `rpbnb_boundary_tests()` reuses the full fit's exact stored draw matrix (zeroing a column) rather than regenerating draws from a count, so passing `boundary_draws` under `engine = "classic"` is an error. For finer control still – testing only "sd" or only "dispersion" (both engines take a `which` argument), or reusing one boundary-test run across several summaries – call `rpbnb_boundary_tests()/rpbnb_tmb_boundary_tests()` directly on the fit and assign its result to `fit$boundary_tests` (with `standardize = TRUE`, reconstruct the fitting-scale data first: `rpbnb:::apply_scaling(data, fit$scaling)`).

Which arguments go with which engine

Argument	engine = "classic"	engine = "tmb"
draw_type, .fixed, .opt_draws	yes	error
inference, keep	error	yes
method, disable_parallel_gaussian, force_parallel_gaussian	ignored with a warning	yes
offset() in a formula	yes	error
dependence = "independence"	error	yes
boundary_draws (non-NULL)	error	yes
control class	rpbnb_control	rpbnb_control
optimizer	maxLik::maxLik(method = "BFGS")	stats::nlmin

Both engines freeze the Famoye admissible lambda interval at the starting values (so the analytic gradient is exactly the derivative of the optimized objective) and validate the fitted lambda against the interval admissible at the optimum afterwards — check `lambda_admissible` on the fit. The fixed-parameter `fit_bnb()` takes the opposite trade-off (moving bounds, admissible by construction); see the decision note in `R/bnb_likelihood.R`.

See Also

`fit_rpbnb()`, `fit_rpbnb_tmb()`, `rpbnb_control()`, `rpbnb_tmb_control()`, `copula()`, `fit_bnb()`

Examples

```
d <- read.csv(system.file("extdata", "rwm1984_bnb.csv", package = "rpbnb"))
fit <- rpbnb(docvis ~ outwork, hospvis ~ outwork, data = d,
            engine = "tmb", random_1 = "outwork", draws = 50)
```

```
# The rpbnb() call from each standalone example under inst/ (see
# inst/example_*.R for the full script: data, formulas, diagnostics).
# Illustration only -- not run by R CMD check.
## Not run:
# example_rpbnb_tmb_famoye.R (sml, Famoye/Sarmanov)
fit_famoye <- rpbnb(f1, f2, data = d, engine = "tmb", method = "sml",
                  random_1 = "kids", random_2 = "kids",
                  draws = DRAWS,
                  seed = SEED,
                  dependence = "famoye",
                  control = ctrl,
                  boundary_tests = LR_TEST)
```

```
# example_rpbnb_tmb_sml.R (sml, copula)
fit_copula <- rpbnb(f1, f2, data = d, engine = "tmb", method = "sml",
                  random_1 = "kids",
                  random_2 = "kids",
                  dependence = copula(COPULA_FAMILY),
                  draws = DRAWS,
                  seed = SEED,
                  control = ctrl,
```

```

        boundary_tests = LR_TEST)

# example_rpbnb_tmb_laplace.R (laplace, copula)
fit_copula <- rpbnb(f1, f2, data = d, engine = "tmb", method = "laplace",
  random_1 = "kids",
  random_2 = "kids",
  dependence = copula(COPULA_FAMILY),
  control = ctrl,
  boundary_tests = LR_TEST)

# example_rpbnb_dense_tmb_sml.R (sml, dense truck-crash data)
fit <- rpbnb(
  formula_1      = f1,
  formula_2      = f2,
  data           = data,
  engine         = "tmb",
  method         = tmb_method,
  boundary_draws = boundary_draws,
  random_1       = random_1,
  random_2       = random_2,
  dependence     = dependence,
  seed           = seed,
  draws          = draws,
  standardize    = TRUE,
  boundary_tests = boundary_tests,
  control        = ctrl
)

# example_rpbnb_dense_tmb_laplace.R (laplace, dense truck-crash data)
fit <- rpbnb(
  formula_1      = f1,
  formula_2      = f2,
  data           = data,
  engine         = "tmb",
  method         = tmb_method,
  random_1       = random_1,
  random_2       = random_2,
  dependence     = dependence,
  standardize    = TRUE,
  boundary_tests = boundary_tests,
  control        = ctrl
)

## End(Not run)

```

Description

Runs a likelihood-ratio test for every random-coefficient standard deviation ($sd1:*$, $sd2:*$) and NB2 dispersion ($m1$, $m2$) of a fitted random-parameter bivariate NB model, and merges them into one table. These are the parameters whose null lies on the boundary of the parameter space ($SD = 0$, or dispersion $m = 0 = \text{Poisson}$), for which the natural-scale summary reports no Wald z/p ; each test uses the 50:50 chi-square boundary correction of `lr_test()` (`boundary = TRUE`).

Usage

```
rpbnb_boundary_tests(
  fit,
  data,
  control = rpbnb_control(compute_se = FALSE),
  which = c("sd", "dispersion")
)
```

Arguments

<code>fit</code>	A converged <code>rpbnb_fit</code> (the full model), from a Famoye or a <code>copula()</code> dependence. Both paths use the exact $m = 0$ branch for the dispersion tests.
<code>data</code>	The data frame the model was fit on. Required – the fit object does not store it, and every restricted model is refit on the same data.
<code>control</code>	An <code>rpbnb_control()</code> for the restricted refits. Defaults to <code>compute_se = FALSE</code> (the LR test needs only <code>logLik</code> and the degrees of freedom). <code>draws/draw_type/seed</code> are taken from <code>fit</code> , not <code>control</code> .
<code>which</code>	Which parameter groups to test, any subset of "sd" (the random-coefficient scales), "dispersion" (the NB2 dispersions $m1$, $m2$), and "dependence" (the association parameter). The default is <code>c("sd", "dispersion")</code> – the two boundary-null groups. "dependence" is opt-in because it costs another full refit and, for three of the four families, tests an <i>interior</i> null whose ordinary Wald z in <code>summary()</code> is already valid. The dependence test restricts the model to independence and compares it to <code>fit</code> : one row labelled <code>lam</code> (Famoye), <code>theta</code> (Frank / Clayton), or <code>rho</code> (Gaussian). The restriction pins the working-scale dependence parameter at its family's independence value rather than refitting a different family, so both fits keep the same draws and the same parameter block and the comparison is a clean 1-df restriction. Only Clayton's null ($\theta > 0$, so $\theta = 0$ is a boundary) takes the 50:50 mixture; Famoye, Frank, and Gaussian nulls are interior and get an ordinary <code>chi-square(1)</code> .

Details

Each test refits a properly nested restricted model. With **multiple random coefficients in an equation, each SD is tested individually** (a 1-df restriction). The restricted fit keeps the full random specification but sets the tested coefficient's draw column to the distribution median ($u = 0.5$), which zeroes that coefficient's per-draw *deviation* exactly for every supported distribution ($u = 0.5$ maps to base 0 for normal/lognormal/ triangular, and to the centered value $2 * 0.5 - 1 = 0$ for uniform). The coefficient therefore collapses to its SD-zero null – the ordinary fixed coefficient b for

normal/uniform/triangular, and $\text{sign} * \exp(b)$ for lognormal – on every draw, independent of the covariate scale. Its (now inert) log-scale is pinned only to drop it from the free-parameter count. Every other draw column is the full model's exact stored draw, so the two simulated log-likelihoods are compared on common random numbers, and each restricted fit is **warm-started from the full fit's coefficients** so the start-sensitive simulated objective does not settle at an inferior optimum. A restricted fit that fails to converge yields NA inference (with a warning) rather than a p-value, and a non-converged full fit is rejected.

Value

An object of class `rpbnb_boundary_tests`: a data frame with columns `Parameter`, `LR`, `df`, `p.value`, `Signif` (one row per boundary parameter), and a `print` method.

Under Famoye dependence with uniform or triangular random coefficients (or a single varying margin), the full and restricted fits' admissible lambda intervals are frozen at different starting values, so each LR statistic compares maxima over slightly different lambda ranges; see the "Famoye caveat" section of `lr_test()`. Normal/lognormal coefficients in both margins are unaffected (the interval is the constant `c(-1, 1)`).

See Also

`lr_test()`, `fit_rpbnb()`

Examples

```
sim <- simulate_rpbnb(n = 600,
  beta1 = c("(Intercept)" = 0.2, x1 = 0.4),
  beta2 = c("(Intercept)" = 0.1, x1 = -0.3),
  random_1 = list(x1 = list(sd = 0.5)),
  dispersion = c(m1 = 0.4, m2 = 0.5), seed = 1)
fit <- fit_rpbnb(y1 ~ x1, y2 ~ x1, data = sim$data, random_1 = "x1",
  draws = 200, seed = 1)
rpbnb_boundary_tests(fit, sim$data)
```

rpbnb_build_info

Report how this installation of rpbnb was compiled

Description

Nearly all of this package's running time is compiled likelihood evaluation, so whether its shared object was built with optimization is a performance fact worth roughly a factor of two on every fit – and nothing in R's own output tells you which build you have. This reports what the compiler actually did.

Usage

`rpbnb_build_info()`

Details

A source install (`install.packages(type = "source")`, R CMD INSTALL, `remotes::install_github()`) compiles with R's own CXXFLAGS, which is `-O2` on every standard platform, so the optimized build is what you get by default. The package does not – and by CRAN policy must not – override those flags. Two things do produce a slow build: a `-O0` (or `-Og`) entry in the user's `~/R/Makevars`, which applies to every package compiled on that machine, and development helpers that inject debug flags of their own, notably `pkgbuild::compile_dll(debug = TRUE)` – its default – which `devtools::load_all()` uses when recompiling changed sources.

Value

A list, invisibly when printed:

`optimized` TRUE when compiled at `-O1` or above, read from the compiler's own `__OPTIMIZE__` macro rather than inferred.

`openmp` TRUE when OpenMP is available, so `n_cores > 1` can do anything. A build without it fits single-threaded whatever `control$n_cores` says.

`openmp_max_threads` The OpenMP runtime's own thread ceiling.

`assertions_enabled` TRUE when `NDEBUG` is unset, a further marker of a debug build.

`compiler` Compiler and version that built the shared object.

Examples

```
rpbnb_build_info()
```

`rpbnb_control`

Control parameters for every rpbnb estimator

Description

One control object for all of the package's fitters – `fit_bnb()`, `fit_rpbnb()`, `fit_rpbnb_tmb()`, and `rpbnb()` with either engine. It carries the union of the tuning knobs the estimators use; each fitter reads the ones that apply to it and **ignores the rest**, reporting the ignored names in `print()/summary()` of the resulting fit rather than erroring. That is what makes a script able to flip `engine = "classic"` to `"tmb"` without rewriting its control call.

Usage

```
rpbnb_control(
  method = c("BFGS"),
  iterlim = NULL,
  reltol = 1e-08,
  print_level = NULL,
  draws_hessian = 100L,
  halton_burn = 300L,
  n_cores = 1L,
```

```

compute_se = TRUE,
hessian = c("numeric", "analytic"),
se_method = NULL,
hess_eps = 1e-05,
hess_r = 4L,
gradtol = 1e-05,
restarts = 10L,
max_threads = NULL,
max_workload = NULL,
parallel_tape = TRUE,
tape_chunks = NULL
)

```

Arguments

method	Optimizer used by the maxLik fitters. Only "BFGS" is implemented and wired through; it is the sole accepted value.
iterlim	Maximum optimizer iterations. NULL (default) uses 300 for the maxLik fitters and 500 for the TMB engine's nlminb.
reltol	Relative convergence tolerance.
print_level	Optimizer verbosity, and the switch that silences the boundary-test progress messages. NULL (default) uses 2 under maxLik and 1 under the TMB engine. Under TMB it drives two separate switches – <code>MakeADFun(silent = print_level == 0)</code> and <code>nlminb's trace = max(0, print_level - 1)</code> – so the levels are: 0 Silent (the pre-0.4.6 TMB default). 1 TMB's own output only: an outer <code>mgc:</code> line per outer evaluation, plus the one-time <code>tape/atomic</code> construction on the first fit of a session. No <code>nlminb</code> trace. The C-level trace lines (<code>Optimizing tape...</code>, <code>Constructing atomic ...</code>, and the <code>N regions found / Using N threads pair</code>) are suppressed when tapes are built concurrently at more than one thread (<code>parallel_tape = TRUE</code>, the default, with <code>n_cores > 1</code>), because TMB emits them from its OpenMP worker threads, where printing aborts the fit; the outer <code>mgc:</code> lines print from the main thread and are unaffected. 2 Adds <code>nlminb</code>'s per-iteration objective and parameter vector – the lowest level that traces every iteration. >2 <code>nlminb</code>'s trace is a print <i>interval</i>, not a verbosity level, so higher values print the objective <i>less</i> often. Note that <code>rpbnb_tmb_boundary_tests()</code> builds its own control when one is not supplied, so its restricted refits print at their own default regardless of this setting.
draws_hessian	Retained for backward compatibility but unused by every estimator: the random-parameter numeric Hessian is taken with the same optimization draws that produced the estimate (same-draw curvature), so it no longer resimulates a separate Hessian draw set. Supplying it is always reported as an ignored setting.
halton_burn	Number of leading Halton points discarded before forming the simulation draws.

n_cores	Worker processes for <code>fit_rpbmb()</code> 's optional cluster path, or OpenMP threads for <code>fit_rpbmb_tmb()</code> (1 = sequential in both cases). Under the TMB engine the <i>realized</i> thread count – after max_threads and hardware capping – is also what max_workload/tape_chunks size the tape against when parallel_tape = TRUE; see "How n_cores, max_workload, and tape_chunks interact" below.
compute_se	If FALSE, skip the Hessian and standard errors. The TMB engine has its own inference argument for this instead.
hessian	How <code>fit_bmb()</code> (famoye) computes the Hessian for standard errors: "numeric" (default, <code>numDeriv::hessian()</code>) or "analytic" (the closed-form Famoye (2010) Appendix Hessian). Both freeze the lambda-bounds at the optimum and yield the same observed-information SEs.
se_method	Standard-error method for <code>fit_rpbmb()</code> (the random-parameter model): "numeric" uses the <code>numDeriv::hessian()</code> observed-information Hessian; "analytic" uses the closed-form observed-information Hessian (Famoye (2010) per-draw second derivatives assembled via the Louis mixture formula) – exact and much faster than "numeric" for larger models; "opg" uses the BHHH / outer-product-of-gradients information from the per-observation scores – fastest (one gradient pass, versus "numeric"'s $O(npar^2)$ finite-difference log-likelihood evaluations), but relies on the information-matrix equality so it is unreliable for parameters at a boundary (e.g. a random-coefficient SD estimated near 0). For copula dependence (<code>fit_rpbmb()</code> with <code>dependence = copula(...)</code>), only "opg" and "numeric" are available; "analytic" is not implemented for the copula path and errors. Default NULL, meaning "this dependence family's own default", resolved once the fit is dispatched (same NULL-until-resolved pattern as <code>iterlim/print_level</code>): "numeric" for <code>dependence = "famoye"</code>, "opg" for <code>dependence = copula(...)</code>. The two differ because their speed/reliability tradeoffs differ – Famoye's fast, exact "analytic" Hessian makes "numeric" (not "opg") the conservative default there, where the copula path has no analytic Hessian at all and its own numeric Hessian is markedly slower than OPG (measured ~3x on an 11-parameter fit) for SEs that mostly agree with it anyway, except (per the boundary caveat above) on random-coefficient scale parameters – fall back to "numeric" for just those, or use <code>rpbmb_boundary_tests()</code>, if precise inference on an SD estimated near 0 matters. Read back a resolved fit's own <code>fit\$control\$se_method</code> (or the value read off <code>object\$control</code> reported by <code>print()/summary()</code>) rather than assuming which default applied.
hess_eps, hess_r	Step and Richardson order for <code>numDeriv::hessian()</code> (used only when the numeric Hessian is selected).
gradtol	Stationarity tolerance for the TMB engine's score, applied <i>relative</i> to the objective: the fit is treated as stationary once <code>max(abs(gradient))</code> falls below <code>gradtol * max(1, abs(nll))</code> . <code>nlminb()</code> declares convergence from its relative function test, which can fire while the score is still far from zero; a Hessian taken there is not a curvature and its inverse can carry negative variances. The fit is restarted until the gradient clears this tolerance; missing it is not warned about on its own, because the copula families legitimately optimise against a clamp or a probability floor where no step improves the objective, but it is reported in the

	warning raised when the Hessian actually fails to factor, and is kept on the fit as <code>optimizer\$max_abs_gradient</code> . The scaling matters because the score is a sum over observations, so an absolute cutoff would tighten with sample size for no statistical reason.
<code>restarts</code>	Maximum number of times to restart <code>nlminb()</code> from its own answer while <code>gradtol</code> is unmet. Restarting resets the trust region and step scaling that stalled the first solve. Set to 0L for the single-call behaviour of earlier versions.
<code>max_threads</code>	Maximum OpenMP threads permitted for one TMB fit. NULL (default) means <code>n_cores</code> , so threads are not capped unless set explicitly below <code>n_cores</code> .
<code>max_workload</code>	Maximum weighted observation-draw evaluations permitted before TMB tape construction; Inf disables the guard. The default and every figure here are derived from TAPE_CALIBRATION, so this text cannot drift from the shipped behaviour. With <code>parallel_tape = FALSE</code> the budget is per fit; with the default <code>parallel_tape = TRUE</code> the tapes are built concurrently and the guard multiplies the workload by the realized thread count (see <code>n_cores</code> above for how that count is realized). One unit is one weighted observation-draw. All figures are measured by <code>inst/dev/tmb_benchmark_memory</code> whose raw results are stored in <code>inst/extdata/memory_calibration.csv</code>. Retained tape size depends on $n * \text{draws}$ alone: $\text{tape (MiB)} = 13.374 + 0.001164 * \text{units}$, with $R^2 = 0.9994$ (1221 bytes per unit over the largest workloads; per-unit cost is higher at small workloads because of the fixed intercept). The budget is set on <i>peak</i> working set, not on retained tape, because peak is what exhausts memory: TMB records the whole likelihood before pruning it to each parallel region, so peak runs about 6.11 times the retained tape plus first-evaluation growth, or 12083 bytes per unit measured directly against peak. The default of 700,000 units therefore targets a peak of about 8 GiB for one fit. Families cost different amounts per unit, so each carries a weight – the largest <i>peak</i> ratio to Famoye observed at matched workload, rounded up to the next tenth: independence 0.7, famoye 1, frank 3.6, gaussian 0.9, clayton 1.1. Frank peaks at over three and a half times Famoye, so treating it as unweighted would under-budget it more than threefold. Raise <code>max_workload</code> deliberately against the memory you actually have, not to make one particular dataset fit. As of <code>rpbnb_tmb_max_workload()</code>, the value <code>rpbnb_tmb_control()</code> actually uses by default is no longer the fixed figure above: it will auto-detect available memory and budget 80% of it, falling back to the fixed 8 GiB figure only when detection is unavailable on the current platform. Call <code>rpbnb_tmb_max_workload</code> directly to set your own budget or detection fraction.
<code>parallel_tape</code>	Construct per-thread TMB tapes concurrently. The default TRUE builds them in parallel for faster tape construction; objective and gradient evaluation remains parallel either way. Set FALSE to build tapes sequentially instead and reduce peak memory. Concurrent tape construction multiplies the per-tape memory workload the <code>max_workload/tape_chunks</code> guard sizes against by the realized thread count (see <code>max_workload</code> above); pairing <code>parallel_tape = TRUE</code> with <code>max_workload = Inf</code> disables that guard entirely; <code>rpbnb_control()</code> warns when it sees that combination. Concurrent construction at a realized count above one

also suppresses TMB's C-level tape, atomic, and parallel-region trace lines at `print_level >= 1` – see `print_level` below.

`tape_chunks` TMB engine, SML fits only. Number of draw chunks to split draws into (see draws at `fit_rpbnb_tmb()`). NULL (default) auto-selects the smallest sufficient count when the weighted workload exceeds `max_workload`, or 1L (no chunking) when it does not. Set explicitly to pin a layout regardless of the auto-threshold; must not exceed draws. Chunking is exact for the requested draws (not an approximation) at the cost of somewhat slower gradient evaluations, and a chunked fit has no taped Hessian – `confint(method = "profile")/rpbnb_tmb_dependence_profile()` fall back to a Wald interval with a warning; Wald/optimHess inference (the default) is unaffected. Ignored for `method = "laplace"` (which has no draw dimension to chunk) and by every non-TMB estimator.

Value

An object of class `c("rpbnb_control", "rpbnb_tmb_control")` (a named list). It carries both class names so that every historical `inherits(control, ...)` check in the package and in user code accepts it.

Which parameters apply to which estimator

Parameter	<code>fit_bnb()</code>	<code>fit_rpbnb()</code>	<code>fit_rpbnb_tmb()</code>
<code>method</code> , <code>compute_se</code> , <code>hess_eps</code> , <code>hess_r</code>	yes	yes	ignored
<code>iterlim</code> , <code>reitol</code> , <code>print_level</code>	yes	yes	yes
<code>halton_burn</code> , <code>n_cores</code>	ignored	yes	yes
<code>hessian</code>	yes	ignored	ignored
<code>se_method</code>	ignored	yes	ignored
<code>gradtol</code> , <code>restarts</code> , <code>max_threads</code> , <code>max_workload</code> , <code>parallel_tape</code>	ignored	ignored	yes
<code>draws_hessian</code>	ignored	ignored	ignored

Note that `iterlim` and `n_cores` mean different things to different estimators – a `maxLik` BFGS iteration limit versus an `nlminb` one, worker *processes* versus OpenMP *threads*. They are not translated; each estimator reads the number and applies its own meaning to it.

Defaults that depend on the estimator

`iterlim` and `print_level` default to NULL, which means "this estimator's own default": `iterlim` is 300 under `maxLik` and 500 under `nlminb`. `print_level` is 2 under `maxLik` and, as of 0.4.6, 1 under the TMB engine (it was 0 – fully silent – before that, which left a long TMB fit printing nothing at all while it ran). The TMB default of 1 shows TMB's own progress without `nlminb`'s per-iteration parameter vectors; see `print_level` below for what each level prints. Supplying either explicitly overrides that for every estimator; `print_level = 0` restores silence. `max_threads` defaults to `n_cores` and `max_workload` is computed from available memory by `rpbnb_tmb_max_workload()` the first time a TMB fit needs it (so a non-TMB fit never pays for the memory probe).

How `n_cores`, `max_workload`, and `tape_chunks` **interact** (TMB, `method = "sml"`)

These three only interact through `parallel_tape`. Before `fit_rpbnb_tmb()` builds a tape, it estimates a weighted workload: `nrow(data) * draws * <family weight> * (realized n_cores if parallel_tape = TRUE)`. `max_workload` caps that number, and `tape_chunks` is what lets a fit exceed it anyway by building several smaller tapes instead of one.

- `n_cores` sets OpenMP threads for the objective/gradient, which are always evaluated in parallel regardless of anything else here. With the default `parallel_tape = TRUE`, the *realized* thread count (after `max_threads` and hardware capping) also multiplies the workload above, because each thread builds its own full tape concurrently – so raising `n_cores` speeds up evaluation but can push a fit that fit comfortably at one thread over `max_workload` at eight. Set `parallel_tape = FALSE` to decouple the two: tapes then build one at a time (peak memory unaffected by `n_cores`) while evaluation stays fully parallel.
- `max_workload` (NULL auto-detects available memory; see above) is checked against that workload number. Exceeding it does not always error: under `method = "sml"` the fit auto-chunks instead (next point). Inf disables the check *and* auto-chunking entirely, so nothing sizes the tape for you – pin `tape_chunks` yourself if you still want the memory benefit with the guard off.
- `tape_chunks` (SML fits only; ignored under `method = "laplace"`, whose tape scales with `nrow(data)` alone, not `draws`) is what absorbs a workload over budget: NULL auto-selects the smallest chunk count that brings the per-chunk workload back under `max_workload`; an explicit value pins the layout (validated against both `draws` and the budget). More chunks trade slower gradient evaluations (each chunk's contribution is recomputed on every outer step) and the loss of a taped Hessian (`confint(method = "profile")/rpbnb_tmb_dependence_profile()` fall back to a Wald interval) for lower peak memory.

Net effect for a large SML fit: pick `n_cores` for speed first. If that trips `max_workload`, either let auto-chunking absorb it, set `parallel_tape = FALSE` if the extra peak memory isn't worth the tape-build speedup, or raise `max_workload` deliberately against memory you actually have (see `rpbnb_tmb_max_workload()`).

See Also

`rpbnb_tmb_max_workload()`, `rpbnb()`, `fit_rpbnb()`, `fit_rpbnb_tmb()`, `fit_bnb()`

Examples

```
rpbnb_control(method = "BFGS", iterlim = 200)
rpbnb_control(hessian = "analytic")
# The same object drives either engine; the TMB-only knobs are simply
# ignored by the classic one (and reported as ignored in its summary).
rpbnb_control(n_cores = 4, gradtol = 1e-6, se_method = "opg")
```

rpbnb_elasticities	<i>Elasticities and semi-elasticities for a random-parameter bivariate NB model</i>
--------------------	---

Description

Continuous elasticities $x_{ij} (\partial \mu_i / \partial x_{ij}) / \mu_i$ and binary semi-elasticities $\mu_i(x_j = 1) / \mu_i(x_j = 0) - 1$, built on the Monte-Carlo integrated population mean $\mu_i = E_\beta[\exp(x'_i \beta)]$ of an `fit_rpbnb()` fit. Under fixed coefficients these reduce to $\beta_j E[x_j]$ and $\exp(\beta_j) - 1$. Standard errors use a numeric delta method over the equation's mean and log-scale parameters.

Usage

```
rpbnb_elasticities(
  fit,
  which = c("y1", "y2", "both"),
  type = c("AME", "MEM"),
  vars = NULL,
  include_intercept = FALSE,
  digits = 4,
  print_output = TRUE,
  n_cores = 1L,
  scaling = NULL,
  log_vars = NULL
)
```

Arguments

<code>fit</code>	An <code>rpbnb_fit</code> object from <code>fit_rpbnb()</code> .
<code>which</code>	Which margin(s): "y1", "y2", or "both".
<code>type</code>	"AME" (average over the sample) or "MEM" (evaluated at the mean row).
<code>vars</code>	Optional variable names or indices to restrict output.
<code>include_intercept</code>	Logical; include the intercept term.
<code>digits</code>	Number of decimal places for printed output.
<code>print_output</code>	Logical; if FALSE, suppress printing.
<code>n_cores</code>	Number of worker processes for the delta-method standard-error jacobian (1 = sequential, the default). When <code>n_cores > 1</code> , the jacobian's independent per-parameter columns are dispatched across a <code>parallel::makeCluster()</code> cluster (one cluster per call, shared across every equation which computes); results are numerically identical to the sequential path. Falls back to sequential with a warning if the <code>parallel</code> package is unavailable.
<code>scaling</code>	Optional named list of <code>c(center =, scale =)</code> pairs, one per covariate that was centred and/or scaled before fitting (e.g. via <code>rpbnb(standardize = TRUE)</code> , whose <code>\$scaling</code> field is exactly this shape), used to report the results in the covariate's

original units. See `rpbnb_marginal_effects()`'s scaling documentation for the full explanation (elasticities and marginal effects share it verbatim).

`log_vars` Optional character vector naming covariates that are ALREADY a log; see `rpbnb_marginal_effects()`'s `log_vars` documentation.

Value

A data frame (single margin, invisibly) or a named list of data frames (both), each with columns Name, Estimate, StdErr, z, p, Signif, var_type.

See Also

`bnb_elasticities()` for fixed-coefficient `bnb_fit` models.

Examples

```
sim <- simulate_rpbnb(n = 400,
  beta1 = c("(Intercept)" = 0.2, x1 = 0.4),
  beta2 = c("(Intercept)" = 0.1, x1 = -0.3),
  random_1 = list(x1 = list(sd = 0.5)),
  dispersion = c(m1 = 0.4, m2 = 0.5), seed = 1)
fit <- fit_rpbnb(y1 ~ x1, y2 ~ x1, data = sim$data, random_1 = "x1",
  draws = 100)
rpbnb_elasticities(fit, which = "both", type = "AME")
```

rpbnb_marginal_effects

Marginal effects for a random-parameter bivariate NB model

Description

Average marginal effects (AME) or marginal effects at the mean (MEM) for each margin of an `fit_rpbnb()` fit, built on the Monte-Carlo integrated population mean $\mu_i = E_\beta[\exp(x_i'\beta)]$ (the same estimand as `predict_rpbnb_fit()`, reusing the fit's stored draws). Continuous effects are $\partial\mu_i/\partial x_{ij} = \text{mean}_r \text{coef}_{rj} \exp(\text{lp}_{ir})$ (the realized coefficient per draw, so random and fixed columns are handled uniformly); binary (0/1) effects are the integrated discrete difference $E[Y|x_j = 1] - E[Y|x_j = 0]$. Standard errors use a numeric delta method over the equation's mean and log-scale parameters.

Usage

```
rpbnb_marginal_effects(
  fit,
  which = c("y1", "y2", "both", "all"),
  type = c("AME", "MEM"),
  vars = NULL,
  include_intercept = FALSE,
  digits = 4,
```

```

    print_output = TRUE,
    n_cores = 1L,
    scaling = NULL,
    log_vars = NULL
  )

```

Arguments

<code>fit</code>	An <code>rpbnb_fit</code> object from <code>fit_rpbnb()</code> .
<code>which</code>	Which margin(s): "y1", "y2", "both", or "all".
<code>type</code>	"AME" (average over the sample) or "MEM" (effect at the mean row).
<code>vars</code>	Optional variable names or indices to restrict output.
<code>include_intercept</code>	Logical; include the intercept term.
<code>digits</code>	Number of decimal places for printed output.
<code>print_output</code>	Logical; if FALSE, suppress printing.
<code>n_cores</code>	Number of worker processes for the delta-method standard-error jacobian (1 = sequential, the default). When <code>n_cores > 1</code> , the jacobian's independent per-parameter columns are dispatched across a <code>parallel::makeCluster()</code> cluster (one cluster per call, shared across every equation which computes); results are numerically identical to the sequential path. Falls back to sequential with a warning if the <code>parallel</code> package is unavailable.
<code>scaling</code>	Optional named list of <code>c(center =, scale =)</code> pairs, one per covariate that was centred and/or scaled before fitting (e.g. via <code>rpbnb(standardize = TRUE)</code> , whose <code>\$scaling</code> field is exactly this shape), used to report the results in the covariate's original units. Columns not named are left alone, so binary indicators and untransformed variables need no entry. Only the reported quantity is rescaled; the fitted design is not rebuilt. That distinction is not cosmetic when a standardized covariate also carries a random coefficient: the random term enters as $x_std * dev$, so substituting $x_raw = x_std * s + c$ would add a $(c/s) * dev$ random intercept the model never estimated – the disguised random intercept centring was meant to remove, back again. The chain rule avoids that: $d(\mu)/d(x_raw) = (d(\mu)/d(x_std))/s$, and the elasticity's leading factor becomes $x_raw/s = x_std + c/s$. Binary effects are untouched.
<code>log_vars</code>	Optional character vector naming covariates that are ALREADY a log, so results are reported per unit of the underlying variable $v_j = \exp(x_j)$ rather than per unit of x_j itself. Without this the elasticity formula treats <code>log v</code> as an ordinary regressor and returns $\bar{x} * b$, the elasticity with respect to the LOG – for a log-traffic column that can be an order of magnitude off from the elasticity with respect to traffic itself (the coefficient). Rejects binary columns (a 0/1 indicator is not a log-transformed variable).

Details

For a model fitted with an `offset()`, absolute marginal effects scale with the offset-inclusive mean: AME uses each observation's training offset, and MEM uses the mean training offset.

Value

A data frame (single margin, invisibly) or a named list of data frames (both/all), each with columns Name, Estimate, StdErr, z, p, Signif, var_type.

See Also

[bnb_marginal_effects\(\)](#) for fixed-coefficient bnb_fit models.

Examples

```
sim <- simulate_rpbnb(n = 400,
  beta1 = c("(Intercept)" = 0.2, x1 = 0.4),
  beta2 = c("(Intercept)" = 0.1, x1 = -0.3),
  random_1 = list(x1 = list(sd = 0.5)),
  dispersion = c(m1 = 0.4, m2 = 0.5), seed = 1)
fit <- fit_rpbnb(y1 ~ x1, y2 ~ x1, data = sim$data, random_1 = "x1",
  draws = 100)
rpbnb_marginal_effects(fit, which = "y1", type = "AME")
```

rpbnb_threads

Number of CPU threads available for the multithreaded likelihood

Description

Number of CPU threads available for the multithreaded likelihood

Usage

```
rpbnb_threads()
```

Value

Integer thread count reported by OpenMP (1 if OpenMP is unavailable).

rpbnb_tmb_boundary_tests

Boundary-corrected LR tests for an rpbnb_tmb_fit's boundary parameters

Description

Runs a likelihood-ratio test for every random-coefficient scale (sd1:*, sd2:*) and NB2 dispersion (m1, m2) of a fitted TMB-engine random-parameter bivariate NB model, and merges them into one table. These are the parameters whose null lies on the boundary of the parameter space (scale = 0, or dispersion m = 0 = Poisson), for which `summary(fit)` reports no Wald z/p; each test uses the 50:50 chi-square boundary correction of [lr_test\(\)](#) (boundary = TRUE).

Usage

```
rpbnb_tmb_boundary_tests(
  fit,
  data,
  control = NULL,
  which = c("sd", "dispersion"),
  draws = fit$draws,
  disable_parallel_gaussian = FALSE,
  force_parallel_gaussian = NULL,
  sml_fallback = TRUE
)
```

Arguments

fit	A converged rpbnb_tmb_fit (from <code>fit_rpbnb_tmb()</code> or <code>rpbnb(engine = "tmb")</code>).
data	The data frame the model was fit on. Required – the fit object does not store it, and every restricted model is refit on the same data. With <code>standardize = TRUE</code> , pass the standardized data (<code>rpbnb(engine = "tmb", boundary_tests = TRUE)</code> does this automatically; a manual call needs <code>rpbnb:::apply_scaling(data, fit\$scaling)</code>).
control	An <code>rpbnb_tmb_control()</code> for the restricted refits. Defaults to <code>rpbnb_tmb_control(print_level = 1, max_workload = fit\$max_workload, tape_chunks = fit\$tape_chunks)</code> – the same <code>n_cores</code> the original fit was called with (1 if fit predates the stored <code>\$parallel</code> field). <code>max_workload</code> propagates fit's own value when fit draw-chunked (see draws at <code>fit_rpbnb_tmb()</code>), so a restricted refit re-chunks to stay in budget instead of rebuilding one full tape; both stay at today's pre-chunking defaults (Inf, NULL) when fit did not chunk (or predates <code>\$tape_integration</code>). When this draws equals <code>fit\$draws</code> (the default), <code>tape_chunks</code> also propagates fit's own resolved chunk count – needed because a PINNED <code>tape_chunks</code> can coexist with <code>max_workload = Inf</code> (propagating budget alone would then give the refit no way to reconstruct that layout, and it would silently build one full tape). A draws that differs from <code>fit\$draws</code> gets a fresh layout resolved from the propagated <code>max_workload</code> alone, never a stale chunk count carried over from a different draw request. <code>seed/method</code> are taken from fit, not control.
which	Which parameter groups to test, any subset of "sd" (the random-coefficient scales), "dispersion" (the NB2 dispersions <code>m1</code> , <code>m2</code>), and "dependence" (the association parameter). The default is <code>c("sd", "dispersion")</code> – the two boundary-null groups. "dependence" is opt-in because it costs another full refit and, for three of the four families, tests an <i>interior</i> null whose ordinary Wald z in <code>summary()</code> is already valid. The dependence test refits the model with <code>dependence = "independence"</code> – the template's own family, which maps <code>z_dep</code> out of the free parameters, giving an exact 1-df restriction on the same draws – and reports one row labelled <code>lam</code> (Famoye), <code>theta</code> (Frank / Kimeldorf), or <code>rho</code> (Gaussian). Only Kimeldorf's null (<code>theta > 0</code> , so <code>theta = 0</code> is a boundary) takes the 50:50 mixture; Famoye, Frank, and Gaussian nulls are interior and get an ordinary chi-square(1).
draws	Number of Halton simulation draws for the restricted refits. Defaults to <code>fit\$draws</code> – the same number fit itself used. Raising this trades refit cost for precision

without needing to refit `fit` at a higher draws; lowering it is cheaper but noisier. A value other than `fit$draws` still shares `fit$seed` (so the Halton sequence's *prefix* is identical) but no longer draws the exact same simulated log-likelihood surface as `fit`, so the LR statistic picks up simulation noise beyond the restriction under test – prefer the default unless you have a specific reason to diverge.

`disable_parallel_gaussian`

Opt-out that restricts each restricted refit's Gaussian-copula evaluation to one thread (see `?fit_rpbnb_tmb`), forwarded to every refit. Default FALSE (multi-threaded, the 0.4.6+ default). This is intentionally a separate argument rather than something read off `fit`: `fit` does not record whether the original fit opted out, so a caller who wants single-threaded refits must pass `disable_parallel_gaussian = TRUE` here even when the original `fit_rpbnb_tmb()/rpbnb()` call also did.

`force_parallel_gaussian`

Deprecated (pre-0.4.6 polarity); use `disable_parallel_gaussian` instead. TRUE (the old opt-in) is a no-op beyond its deprecation warning; an explicit FALSE is honored as `disable_parallel_gaussian = TRUE`.

`sml_fallback`

When `fit` was estimated by `method = "laplace"` and a restricted refit's Laplace pair cannot be trusted, re-run **that one test's** LR with both sides estimated by `method = "sml"` instead of reporting NA (or a clamped 0). Default TRUE. Two triggers: the restricted refit **fails to converge**, or it reports convergence with a restricted `logLik` *above* the full fit's (a negative raw LR). The second is not a rounding curiosity: the models are nested, so at true optima the restricted `logLik` can never exceed the full fit's – a negative LR proves at least one Laplace value is wrong, either a full fit that stopped short of its optimum (observed at -0.002 to -1.2 nats: the warm-started restricted refit out-polished it) or a restricted fit that climbed a spurious ridge of the approximation itself (observed at -3838 nats: near a singular inner Hessian the Laplace log-likelihood rises without bound, and `nlminb` reports code 0 there). Clamping such a statistic to 0 would silently turn a real effect into "no evidence".

This exists because some restrictions leave Laplace no valid optimum to converge to. Pinning a margin to Poisson can push the dependence strong enough (observed on the truck data's `m1` test under a Frank copula: `theta` driven to about 20, Kendall's `tau` 0.8) that the per-observation cell probability is non-log-concave in the random effects – and the Laplace approximation differentiates through an inner Newton that requires exactly the log-concavity the restricted model no longer has. No inner-solver setting fixes that (`tol10 = 0` clears TMB's "Newton drop out" but the outer fit still ends at `nlminb` false convergence (8) with a gradient of `1e10`); SML has no inner Newton and is not exposed to it.

The fallback refits the **full model too** under SML (once, cached across tests) at these same draws and `fit$seed`, and takes the LR between the two SML fits – never between a Laplace `logLik` and an SML `logLik`, which are different approximations of the likelihood and whose difference is not an LR statistic. Rows that used the fallback are listed in the result's `sml_fallback` attribute and announced by a `message()` (silenced by `control$print_level = 0` like the other announcements). If the SML pair fails to converge as well, the row is NA with a warning, exactly as before. Ignored when `fit` itself was estimated by SML (there is nothing different to fall back to).

Details

Each test refits a properly nested restricted model, otherwise identical to `fit` (same formulas, random-coefficient specification, dependence, seed, and estimator, and by default the same draws), warm-started from `fit$coef` and run with `inference = "none"` since the LR test needs only `logLik` and the parameter count.

Dispersions are restricted with `poisson_1/poisson_2 = TRUE`, the template's exact `m = 0` branch.

Scales are restricted by pinning that coefficient's `log_sd` at the parameterization's zero (`-20`; the template clamps `log_sd` to `[-20, 20]` and computes `sd = exp(log_sd)`, so this is `sd = 2.1e-9` – numerically zero on every draw) and holding it out of the free-parameter count, giving a 1-df restriction. With **multiple random coefficients in an equation, each scale is tested individually**.

Pinning the scale rather than dropping the coefficient from `random_1/random_2` is what preserves **common random numbers**: the Halton draw matrix keeps the same width, so every *other* random coefficient draws from exactly the dimensions it did in the full fit, and the two simulated log-likelihoods differ only by the restriction under test. Dropping the name instead would renumber the remaining coefficients' Halton dimensions, and the LR statistic would absorb that reshuffling as extra simulation noise.

A `message()` ("Boundary LR test: <parameter>...") reports each restricted refit right before it starts, unless `control$print_level` is 0; suppress it with `suppressMessages()` if needed.

Value

An object of class `rpbnb_boundary_tests` – the same class `rpbnb_boundary_tests()` returns, with columns `Parameter`, `LR`, `df`, `p.value`, `Signif` (one row per boundary parameter) and the same `print()` method – so the two engines' results are interchangeable wherever that class is consumed (e.g. `summary.rpbnb_tmb_fit()`'s scale and dispersion blocks, once `$boundary_tests` is attached to the fit).

Scale rows are labelled by the distribution's own scale parameter (`sd` for normal/lognormal, `w` for uniform/triangular half-width, `s` for a lognormal log-scale), matching `summary()`'s row names.

The `sml_fallback` attribute is a character vector of the `Parameter` rows whose LR came from the SML fallback pair (see the `sml_fallback` argument); `character(0)` when every test ran under `fit`'s own estimator.

See Also

`lr_test()`, `fit_rpbnb_tmb()`, `rpbnb_boundary_tests()` (the classic-engine counterpart)

Examples

```
sim <- simulate_rpbnb(n = 600,
  beta1 = c("(Intercept)" = 0.2, x1 = 0.4),
  beta2 = c("(Intercept)" = 0.1, x1 = -0.3),
  random_1 = list(x1 = list(sd = 0.5)),
  dispersion = c(m1 = 0.4, m2 = 0.5), seed = 1)
fit <- fit_rpbnb_tmb(y1 ~ x1, y2 ~ x1, data = sim$data, random_1 = "x1",
  draws = 100)
rpbnb_tmb_boundary_tests(fit, sim$data)
```

rpbnb_tmb_control	<i>Control parameters for the TMB engine (alias of rpbnb_control())</i>
-------------------	---

Description

Retained so that code written against the pre-unification API keeps working. The two control objects were merged in 0.4.1: this function forwards to `rpbnb_control()` and returns exactly the same object, which every estimator in the package accepts. New code should call `rpbnb_control()` directly.

Usage

```
rpbnb_tmb_control(
  iterlim = NULL,
  reltol = 1e-08,
  gradtol = 1e-05,
  restarts = 10L,
  print_level = NULL,
  n_cores = 1L,
  max_threads = NULL,
  max_workload = NULL,
  parallel_tape = TRUE,
  halton_burn = 300L,
  tape_chunks = NULL
)
```

Arguments

<code>iterlim</code>	Maximum optimizer iterations. NULL (default) uses 300 for the maxLik fitters and 500 for the TMB engine's nlminb.
<code>reltol</code>	Relative convergence tolerance.
<code>gradtol</code>	Stationarity tolerance for the TMB engine's score, applied <i>relative</i> to the objective: the fit is treated as stationary once $\max(\text{abs}(\text{gradient}))$ falls below $\text{gradtol} * \max(1, \text{abs}(\text{nl}))$. <code>nlminb()</code> declares convergence from its relative function test, which can fire while the score is still far from zero; a Hessian taken there is not a curvature and its inverse can carry negative variances. The fit is restarted until the gradient clears this tolerance; missing it is not warned about on its own, because the copula families legitimately optimise against a clamp or a probability floor where no step improves the objective, but it is reported in the warning raised when the Hessian actually fails to factor, and is kept on the fit as <code>optimizer\$max_abs_gradient</code> . The scaling matters because the score is a sum over observations, so an absolute cutoff would tighten with sample size for no statistical reason.
<code>restarts</code>	Maximum number of times to restart <code>nlminb()</code> from its own answer while <code>gradtol</code> is unmet. Restarting resets the trust region and step scaling that stalled the first solve. Set to 0L for the single-call behaviour of earlier versions.

- `print_level` Optimizer verbosity, and the switch that silences the boundary-test progress messages. NULL (default) uses 2 under `maxLik` and 1 under the TMB engine. Under TMB it drives two separate switches – `MakeADFun(silent = print_level == 0)` and `nlminb's trace = max(0, print_level - 1)` – so the levels are:
- 0 Silent (the pre-0.4.6 TMB default).
 - 1 TMB's own output only: an outer `mgc`: line per outer evaluation, plus the one-time tape/atomic construction on the first fit of a session. No `nlminb` trace. The C-level trace lines (Optimizing tape... , Constructing atomic ..., and the N regions found / Using N threads pair) are suppressed when tapes are built concurrently at more than one thread (`parallel_tape = TRUE`, the default, with `n_cores > 1`), because TMB emits them from its OpenMP worker threads, where printing aborts the fit; the outer `mgc`: lines print from the main thread and are unaffected.
 - 2 Adds `nlminb`'s per-iteration objective and parameter vector – the lowest level that traces every iteration.
 - >2 `nlminb`'s trace is a print *interval*, not a verbosity level, so higher values print the objective *less* often.

Note that `rpbmb_tmb_boundary_tests()` builds its own control when one is not supplied, so its restricted refits print at their own default regardless of this setting.

`n_cores` Worker processes for `fit_rpbmb()`'s optional cluster path, or OpenMP threads for `fit_rpbmb_tmb()` (1 = sequential in both cases). Under the TMB engine the *realized* thread count – after `max_threads` and hardware capping – is also what `max_workload/tape_chunks` size the tape against when `parallel_tape = TRUE`; see "How `n_cores`, `max_workload`, and `tape_chunks` interact" below.

`max_threads` Maximum OpenMP threads permitted for one TMB fit. NULL (default) means `n_cores`, so threads are not capped unless set explicitly below `n_cores`.

`max_workload` Maximum weighted observation-draw evaluations permitted before TMB tape construction; Inf disables the guard. The default and every figure here are derived from `TAPE_CALIBRATION`, so this text cannot drift from the shipped behaviour.

With `parallel_tape = FALSE` the budget is per fit; with the default `parallel_tape = TRUE` the tapes are built concurrently and the guard multiplies the workload by the realized thread count (see `n_cores` above for how that count is realized).

One unit is one weighted observation-draw. All figures are measured by `inst/dev/tmb_benchmark_memory` whose raw results are stored in `inst/extdata/memory_calibration.csv`.

Retained tape size depends on $n * \text{draws}$ alone: $\text{tape (MiB)} = 13.374 + 0.001164 * \text{units}$, with $R^2 = 0.9994$ (1221 bytes per unit over the largest workloads; per-unit cost is higher at small workloads because of the fixed intercept).

The budget is set on *peak* working set, not on retained tape, because peak is what exhausts memory: TMB records the whole likelihood before pruning it to each parallel region, so peak runs about 6.11 times the retained tape plus first-evaluation growth, or 12083 bytes per unit measured directly against peak. The default of 700,000 units therefore targets a peak of about 8 GiB for one fit.

Families cost different amounts per unit, so each carries a weight – the largest *peak* ratio to Famoye observed at matched workload, rounded up to the next

tenth: independence 0.7, famoye 1, frank 3.6, gaussian 0.9, clayton 1.1. Frank peaks at over three and a half times Famoye, so treating it as unweighted would under-budget it more than threefold.

Raise `max_workload` deliberately against the memory you actually have, not to make one particular dataset fit.

As of `rpbnb_tmb_max_workload()`, the value `rpbnb_tmb_control()` actually uses by default is no longer the fixed figure above: it will auto-detect available memory and budget 80% of it, falling back to the fixed 8 GiB figure only when detection is unavailable on the current platform. Call `rpbnb_tmb_max_workload` directly to set your own budget or detection fraction.

<code>parallel_tape</code>	Construct per-thread TMB tapes concurrently. The default TRUE builds them in parallel for faster tape construction; objective and gradient evaluation remains parallel either way. Set FALSE to build tapes sequentially instead and reduce peak memory. Concurrent tape construction multiplies the per-tape memory workload the <code>max_workload/tape_chunks</code> guard sizes against by the realized thread count (see <code>max_workload</code> above); pairing <code>parallel_tape = TRUE</code> with <code>max_workload = Inf</code> disables that guard entirely; <code>rpbnb_control()</code> warns when it sees that combination. Concurrent construction at a realized count above one also suppresses TMB's C-level tape, atomic, and parallel-region trace lines at <code>print_level >= 1</code> – see <code>print_level</code> below.
<code>halton_burn</code>	Number of leading Halton points discarded before forming the simulation draws.
<code>tape_chunks</code>	TMB engine, SML fits only. Number of draw chunks to split draws into (see draws at <code>fit_rpbnb_tmb()</code>). NULL (default) auto-selects the smallest sufficient count when the weighted workload exceeds <code>max_workload</code> , or 1L (no chunking) when it does not. Set explicitly to pin a layout regardless of the auto-threshold; must not exceed draws. Chunking is exact for the requested draws (not an approximation) at the cost of somewhat slower gradient evaluations, and a chunked fit has no taped Hessian – <code>confint(method = "profile")/rpbnb_tmb_dependence_profile()</code> fall back to a Wald interval with a warning; Wald/optimHess inference (the default) is unaffected. Ignored for <code>method = "laplace"</code> (which has no draw dimension to chunk) and by every non-TMB estimator.

Details

Only the arguments you actually supply are forwarded, so an untouched `iterlim/print_level` still resolves to the TMB engine's own defaults (500 and 1) when the object is used for a TMB fit – and to the `maxLik` defaults if the same object is handed to `fit_rpbnb()`.

Value

The `rpbnb_control()` object.

See Also

`rpbnb_control()`

Examples

```
identical(unclass(rpbnb_tmb_control(n_cores = 2)),
          unclass(rpbnb_control(n_cores = 2)))
```

```
rpbnb_tmb_dependence_profile
```

Confidence interval for a fitted dependence parameter

Description

Returns a profile-likelihood interval for the dependence parameter of a fitted model, computed on the unconstrained working scale and mapped through the family's monotone link. This is the tool to reach for when `summary()` reports NA for the dependence standard error: at a boundary the delta-method derivative collapses, so $SE = |d\theta/dz| \cdot SE(z)$ is a $0 \times \infty$ product and no symmetric standard error exists. A profile interval needs no derivative.

Usage

```
rpbnb_tmb_dependence_profile(
  fit,
  level = 0.95,
  method = c("profile", "wald"),
  ...
)
```

Arguments

<code>fit</code>	An object of class <code>rpbnb_tmb_fit</code> .
<code>level</code>	Coverage level; one number strictly between 0 and 1.
<code>method</code>	"profile" (default) for a profile-likelihood interval, or "wald" for a Wald interval on the working scale mapped through the same link. "profile" requires the TMB objective, retained only under <code>keep = "full"</code> ; when it is unavailable this degrades to "wald" with a warning rather than failing.
<code>...</code>	Passed to <code>tmbprofile</code> , e.g. <code>ytol</code> or <code>parm.range</code> . <code>lincomb</code> and <code>slice</code> are not supported and raise an error: the first would profile a different quantity than <code>z_dep</code> while still being reported as if it were, and the second returns a likelihood slice rather than a profile.

Details

For Famoye this is the usual situation, because the admissible lambda interval is frozen at the starting values (see `lambda_bounds` in `fit_rpbnb_tmb`). An interval around a pinned estimate is still an interval around an artefact – widen the box first by refitting from better starting values.

Value

A data frame with one row per reported dependence quantity and columns parameter, estimate, lower, upper, level, and method – the method actually used, which may differ from the one requested. The raw profile, when one was computed, is attached as `attr(, "profile")` so it can be plotted.

See Also

[fit_rpbnb_tmb](#)

Examples

```
## Not run:
fit <- fit_rpbnb_tmb(y1 ~ x, y2 ~ x, data = d,
                    dependence = "famoye", keep = "full")
rpbnb_tmb_dependence_profile(fit)
plot(attr(rpbnb_tmb_dependence_profile(fit), "profile"))

## End(Not run)
```

rpbnb_tmb_elasticities

Elasticities for a rpbnb_tmb model

Description

Continuous elasticities $x_j / E[Y] \cdot \partial E[Y] / \partial x_j$ and binary semi-elasticities $E[Y|x_j = 1] / E[Y|x_j = 0] - 1$. Built on the same Monte-Carlo integrated mean used by [rpbnb_tmb_marginal_effects](#).

Usage

```
rpbnb_tmb_elasticities(
  fit,
  which = c("y1", "y2", "both"),
  type = c("AME", "MEM"),
  vars = NULL,
  include_intercept = FALSE,
  digits = 4L,
  scaling = NULL,
  log_vars = NULL,
  ...
)
```

Arguments

<code>fit</code>	An object of class <code>rpbnb_tmb_fit</code> .
<code>which</code>	Which margin: "y1", "y2", or "both".
<code>type</code>	"AME" (average over the sample) or "MEM" (at the sample mean of covariates).
<code>vars</code>	Optional character vector of variable names to restrict output.
<code>include_intercept</code>	Logical; include the intercept term in output.
<code>digits</code>	Number of decimal places for printed table entries.
<code>scaling</code>	Optional named list of <code>c(center =, scale =)</code> pairs, one per covariate that was centred and/or scaled before fitting, used to report the results in the covariate's original units. Columns not named are left alone, so binary indicators and untransformed variables need no entry. Only the reported quantity is rescaled; the fitted design is not rebuilt. That distinction is not cosmetic when a standardized covariate also carries a random coefficient. The random term enters as $x_{std} * dev$, so substituting $x_{raw} = x_{std} * s + c$ would add a $(c/s) * dev$ random intercept the model never estimated – which is how a centred random slope turns back into the disguised random intercept that centring was meant to remove. The chain rule avoids that entirely: $\partial \mu / \partial x_{raw} = (\partial \mu / \partial x_{std}) / s$, and the elasticity's leading factor becomes $x_{raw} / s = x_{std} + c / s$. Binary effects are untouched.
<code>log_vars</code>	Optional character vector naming covariates that are ALREADY a log, so that results are reported per unit of the underlying variable $v_j = \exp(x_j)$ rather than per unit of x_j itself. Without this the elasticity formula treats $\log v$ as an ordinary regressor and returns $\bar{x}b$, the elasticity with respect to the LOG – for the truck model's LNAADT_3 that is 8.59, where the elasticity with respect to AADT is the coefficient itself, 0.871. A ten-fold overstatement that looks perfectly plausible in a table. Because $\partial \log \mu / \partial \log v = (\partial \mu / \partial x) / \mu$, the elasticity is $E[d\mu/\mu]/s$ and the marginal effect is $E[d\mu/(sv)]$, with $v = \exp(x_{std}s + c)$ picking up any scaling applied to the logged column. Named columns are reported with type "log-continuous". Binary columns are rejected.
<code>...</code>	Not used.

Value

A data frame or named list of data frames.

`rpbnb_tmb_marginal_effects`

Marginal effects for a rpbnb_tmb model

Description

Computes average marginal effects (AME) or marginal effects at the mean (MEM) for a fitted random-parameter bivariate negative binomial model. For continuous covariates the effect is $\partial E[Y]/\partial x_j$; for binary covariates it is the discrete difference $E[Y|x_j = 1] - E[Y|x_j = 0]$. When a coefficient is random the per-draw realized coefficients are used to form the Monte-Carlo integrated mean.

Usage

```
rpbnb_tmb_marginal_effects(
  fit,
  which = c("y1", "y2", "both"),
  type = c("AME", "MEM"),
  vars = NULL,
  include_intercept = FALSE,
  digits = 4L,
  scaling = NULL,
  log_vars = NULL,
  ...
)
```

Arguments

fit	An object of class <code>rpbnb_tmb_fit</code> .
which	Which margin: "y1", "y2", or "both".
type	"AME" (average over the sample) or "MEM" (at the sample mean of covariates).
vars	Optional character vector of variable names to restrict output.
include_intercept	Logical; include the intercept term in output.
digits	Number of decimal places for printed table entries.
scaling	Optional named list of <code>c(center =, scale =)</code> pairs, one per covariate that was centred and/or scaled before fitting, used to report the results in the covariate's original units. Columns not named are left alone, so binary indicators and untransformed variables need no entry. Only the reported quantity is rescaled; the fitted design is not rebuilt. That distinction is not cosmetic when a standardized covariate also carries a random coefficient. The random term enters as $x_{std} * dev$, so substituting $x_{raw} = x_{std} * s + c$ would add a $(c/s) * dev$ random intercept the model never estimated – which is how a centred random slope turns back into the disguised random intercept that centring was meant to remove. The chain rule avoids that entirely: $\partial\mu/\partial x_{raw} = (\partial\mu/\partial x_{std})/s$, and the elasticity's leading factor becomes $x_{raw}/s = x_{std} + c/s$. Binary effects are untouched.
log_vars	Optional character vector naming covariates that are ALREADY a log, so that results are reported per unit of the underlying variable $v_j = \exp(x_j)$ rather than per unit of x_j itself. Without this the elasticity formula treats $\log v$ as an ordinary regressor and returns $\bar{x}b$, the elasticity with respect to the LOG – for the truck model's LNAADT_3

that is 8.59, where the elasticity with respect to AADT is the coefficient itself, 0.871. A ten-fold overstatement that looks perfectly plausible in a table.

Because $\partial \log \mu / \partial \log v = (\partial \mu / \partial x) / \mu$, the elasticity is $E[d\mu/\mu]/s$ and the marginal effect is $E[d\mu/(sv)]$, with $v = \exp(x_{std}s + c)$ picking up any scaling applied to the logged column. Named columns are reported with type "log-continuous". Binary columns are rejected.

... Not used.

Value

A data frame (single margin) or named list of two data frames ("both").

rpbnb_tmb_max_workload

Compute a TMB workload budget from a memory figure

Description

Companion to `rpbnb_tmb_control()`'s `max_workload`: rather than picking a weighted-observation-draw count directly, state a memory budget and let this function do the arithmetic TAPE_CALIBRATION implies.

Usage

```
rpbnb_tmb_max_workload(budget_gib = NULL, fraction = 0.8)
```

Arguments

<code>budget_gib</code>	Optional memory budget in GiB, stated explicitly. When supplied, used as-is – fraction does not apply, because a number you state yourself is not second-guessed with a discount. When omitted (the default), available memory is auto-detected and fraction of it is budgeted instead.
<code>fraction</code>	Of auto-detected available memory, the fraction to actually budget; one number in $(0, 1]$. Available memory fluctuates and competes with other processes, so budgeting all of it risks the guard passing a fit that then exhausts memory anyway. Ignored when <code>budget_gib</code> is supplied.

Details

With `budget_gib` omitted, this is also `rpbnb_tmb_control()`'s own default: every fit that doesn't set `max_workload` explicitly already goes through this function.

Value

One positive numeric workload value, on the same scale as `rpbnb_tmb_control()`'s `max_workload`.

Examples

```

rpbnb_tmb_max_workload(budget_gib = 16)
## Not run:
ctrl <- rpbnb_tmb_control(max_workload = rpbnb_tmb_max_workload())

## End(Not run)

```

simulate_bnb

*Simulate data from the Famoye/Sarmanov bivariate NB distribution***Description**

Generates paired count outcomes (y_1 , y_2) from the fixed-parameter Famoye/Sarmanov bivariate NB2 joint PMF:

Usage

```

simulate_bnb(
  n,
  beta1,
  beta2,
  dispersion = c(m1 = 0.5, m2 = 0.5),
  lambda = 0,
  covariates = NULL,
  seed = NULL
)

```

Arguments

<code>n</code>	Number of observations.
<code>beta1, beta2</code>	Named numeric vectors of coefficients for each equation; must include "(Intercept)".
<code>dispersion</code>	Named numeric <code>c(m1 = ..., m2 = ...)</code> NB2 dispersion parameters (variance = $\mu + m * \mu^2$).
<code>lambda</code>	Famoye/Sarmanov dependence parameter. Must lie within the valid bounds implied by the marginal means and dispersions.
<code>covariates</code>	Optional data frame of covariates. If NULL, standard- normal columns are generated for every non-intercept coefficient name.
<code>seed</code>	Optional random seed. If NULL (default) the RNG is left untouched and draws continue from the caller's current stream, so repeated calls yield distinct datasets; supply an integer for reproducible output.

Details

$$P(Y_1 = y_1, Y_2 = y_2) = p_1(y_1) p_2(y_2) [1 + \lambda(e^{-y_1} - c_1)(e^{-y_2} - c_2)]$$

where $c_k = E[e^{-Y_k}]$ under $NB2(\mu_k, m_k)$.

Value

A list with:

`data` data frame with `y1`, `y2`, and covariate columns

`mu` data frame with `mu1`, `mu2` (per-obs conditional means)

`true` list of true parameters: `beta1`, `beta2`, `dispersion`, `lambda`

`settings` list with `n` and `seed`

`meta` list with `seed` and `r_version`

Examples

```
sim <- simulate_bnb(n = 500,
  beta1 = c("(Intercept)" = 0.5, x1 = 0.3),
  beta2 = c("(Intercept)" = 0.2, x1 = -0.2),
  dispersion = c(m1 = 0.4, m2 = 0.5), lambda = 0.1, seed = 1)
head(sim$data)
```

simulate_rpbnb

Simulate data from a random-parameter bivariate NB process

Description

Simulate data from a random-parameter bivariate NB process

Usage

```
simulate_rpbnb(
  n,
  beta1,
  beta2,
  random_1 = NULL,
  random_2 = NULL,
  dispersion = c(m1 = 0.5, m2 = 0.5),
  lambda = 0,
  covariates = NULL,
  seed = NULL
)
```

Arguments

`n` Number of observations.

`beta1`, `beta2` Named numeric vectors of fixed coefficient means per equation; must include "(Intercept)".

random_1, random_2	Named lists giving random coefficients. Each value is a list with dist (one of "normal", "lognormal", "uniform", "triangular"; default "normal"), scale (or sd) for the dispersion, and sign (-1/1, lognormal only). Means come from beta1/beta2; for a lognormal coefficient the beta entry is the log-location and the realized coefficient is $\text{sign} * \exp(\log_location + \text{scale} * z)$.
dispersion	Named numeric $c(m1 = \dots, m2 = \dots)$ NB2 dispersions.
lambda	Famoye dependence parameter (0 = independent margins).
covariates	Optional data frame of covariates; if NULL, standard-normal columns are generated for every non-intercept name.
seed	Optional random seed. If NULL (default) the RNG is left untouched and draws continue from the caller's current stream, so repeated calls yield distinct datasets; supply an integer for reproducible output.

Value

A list with data (y1, y2, covariates), coef_realized (per-obs coefficients per equation), mu (conditional means), true (parameters), settings, and meta (R/seed/timestamp passed by caller).

Examples

```
sim <- simulate_rpbnb(n = 500,
  beta1 = c("(Intercept)" = 0.2, x1 = 0.4),
  beta2 = c("(Intercept)" = 0.1, x1 = -0.3),
  random_1 = list(x1 = list(sd = 0.5)),
  dispersion = c(m1 = 0.4, m2 = 0.5), seed = 1)
head(sim$data)
```

simulate_rpbnb_copula *Simulate data from a copula RP-BNB process*

Description

Simulate data from a copula RP-BNB process

Usage

```
simulate_rpbnb_copula(
  n,
  beta1,
  beta2,
  random_1 = NULL,
  random_2 = NULL,
  dispersion = c(m1 = 0.5, m2 = 0.5),
  copula,
  covariates = NULL,
  seed = NULL
)
```

Arguments

n	Number of observations.
beta1, beta2	Named coefficient means; must include "(Intercept)".
random_1, random_2	Random-coefficient specs (see simulate_rpbnb()).
dispersion	Named c(m1=, m2=) NB2 dispersions.
copula	An copula() object giving the family and native parameter par.
covariates	Optional covariate data frame; NULL -> standard-normal columns.
seed	Optional RNG seed.

Value

list(data, mu, true, settings).

simulate_rpbnb_tmb	<i>Simulate data from a bivariate NB process</i>
--------------------	--

Description

Generates count data from a bivariate negative binomial model with random coefficients and Famoye/Sarmanov or copula dependence.

Usage

```
simulate_rpbnb_tmb(
  n,
  beta1,
  beta2,
  random_1 = NULL,
  random_2 = NULL,
  dispersion = c(m1 = 0.5, m2 = 0.5),
  dependence = "famoye",
  lambda = 0,
  covariates = NULL,
  seed = NULL
)
```

Arguments

n	Number of observations.
beta1, beta2	Named coefficient vectors (must include "(Intercept)").
random_1, random_2	Random coefficient specs (same format as fit_rpbnb_tmb).
dispersion	Named vector c(m1 = , m2 =) of NB2 dispersions.

dependence	"famoye", "independence", or a copula() object.
lambda	Famoye dependence parameter (0 = independence).
covariates	Optional data frame. If NULL, standard-normal covariates generated.
seed	Optional integer seed.

Value

List with \$data (data.frame), \$truth (parameters), \$settings.

Examples

```
sim <- simulate_rpbnb_tmb(n = 200,
  beta1 = c("(Intercept)" = 0.2, x1 = 0.4),
  beta2 = c("(Intercept)" = 0.1, x1 = -0.3),
  dispersion = c(m1 = 0.4, m2 = 0.5), seed = 1)
head(sim$data)
```

summary.rpbnb_tmb_fit *Summarize a fitted TMB-engine rpbnb model*

Description

Summarize a fitted TMB-engine rpbnb model

Usage

```
## S3 method for class 'rpbnb_tmb_fit'
summary(object, digits = 4L, ...)
```

Arguments

object	A fitted model object of class rpbnb_tmb_fit.
digits	Number of decimal places for numeric columns in coefficient tables. Default 4. Use a negative value (e.g. -1) for full precision.
...	Not used.

Value

The fitted object, invisibly.

Index

bnb_boundary_tests, 3
bnb_boundary_tests(), 9
bnb_elasticities, 4
bnb_elasticities(), 38
bnb_gof, 5
bnb_marginal_effects, 5
bnb_marginal_effects(), 40
bnb_residual_checks, 7

copula, 7
copula(), 3, 9, 11, 23, 26, 27, 29, 55

fit_bnb, 8
fit_bnb(), 3–8, 14, 16, 18, 20, 21, 23, 27, 31, 33, 35, 36
fit_rpbnnb, 10, 15
fit_rpbnnb(), 7, 8, 14, 16, 19, 21–23, 27, 30, 31, 33, 35–39, 45, 46
fit_rpbnnb_tmb, 12, 47, 48
fit_rpbnnb_tmb(), 22–24, 26, 27, 31, 33, 35, 36, 41, 43, 45, 46

lr_test, 16
lr_test(), 3, 9, 11, 29, 30, 40, 43

message(), 26, 42, 43

numDeriv::hessian(), 33

plot.bnb_fit, 18
plot.bnb_fit(), 19
plot.rpbnnb_fit, 19
predict.rpbnnb_tmb_fit, 19

residuals.bnb_fit, 20
residuals.rpbnnb_fit, 21
rpbnnb, 22
rpbnnb(), 31, 36
rpbnnb_boundary_tests, 28
rpbnnb_boundary_tests(), 3, 11, 12, 23, 25, 26, 33, 43

rpbnnb_build_info, 30
rpbnnb_control, 31
rpbnnb_control(), 3, 9, 11, 13, 14, 23, 24, 27, 29, 44, 46
rpbnnb_elasticities, 37
rpbnnb_marginal_effects, 38
rpbnnb_marginal_effects(), 38
rpbnnb_threads, 40
rpbnnb_tmb_boundary_tests, 15, 40
rpbnnb_tmb_boundary_tests(), 23, 26
rpbnnb_tmb_control, 44
rpbnnb_tmb_control(), 27, 41
rpbnnb_tmb_dependence_profile, 16, 47
rpbnnb_tmb_elasticities, 48
rpbnnb_tmb_marginal_effects, 48, 49
rpbnnb_tmb_max_workload, 34, 46, 51
rpbnnb_tmb_max_workload(), 35, 36

simulate_bnb, 52
simulate_rpbnnb, 53
simulate_rpbnnb(), 55
simulate_rpbnnb_copula, 54
simulate_rpbnnb_copula(), 7, 8
simulate_rpbnnb_tmb, 55
summary.rpbnnb_tmb_fit, 56
suppressMessages(), 26, 43

tmbprofile, 47